INTELLIGENT PRODUCTION MANAGEMENT SYSTEM

A Project report submitted in partial fulfillment of requirements for the award of degree of

BACHELOR OF TECHNOLOGY
IN
COMPUTER SCIENCE AND ENGINEERING

By

Konda Bhavana (229X1A0514)

Anshu Bhandekar (229X1A0598)

Borelli Tejaswini (229X1A0570)

Under the esteemed guidance of

Sri. B. Dilip Kumar Reddy
Assistant Professor
Department of C.S.E.

Department of Computer Science and Engineering

G. PULLA REDDY ENGINEERING COLLEGE (Autonomous): KURNOOL

(Affiliated to JNTUA, ANANTAPURAMU)

2025 – 2026

Department of Computer Science and Engineering

G. PULLA REDDY ENGINEERING COLLEGE (Autonomous): KURNOOL

(Affiliated to JNTUA, ANANTAPURAMU)

CERTIFICATE

This is to certify that the Project Work entitled 'Intelligent Production Management System' is a bonafide record of work carried out by

Konda Bhavana (229X1A0514)

Anshu Bhandekar (229X1A0598)

Borelli Tejaswini (229X1A0570)

Under my guidance and supervision in partial fulfillment of the requirements for the award of degree of

BACHELOR OF TECHNOLOGY
IN
COMPUTER SCIENCE AND ENGINEERING

……..	……..
Sri. B. Dilip Kumar Reddy	**Dr. N. Kasiviswanath**
Assistant Professor,	Professor & Head of the Department,
Department of C.S.E.,	Department of C.S.E.,
G. Pulla Reddy Engineering College, Kurnool.	G. Pulla Reddy Engineering College, Kurnool.

DECLARATION

We hereby declare that the project titled "**Intelligent Production Management System"** is an authentic work carried out by us as the students of **G. PULLA REDDY ENGINEERING COLLEGE(Autonomous) Kurnool,** during 2025-26 and has not been submitted elsewhere for the award of any degree or diploma in part or in full to any institute.

Konda Bhavana
(229X1A0514)

Anshu Bhandekar
(229X1A0598)

Borelli Tejaswini
(229X1A0570)

ACKNOWLEDGEMENT

We wish to express our deep sense of gratitude to our project guide **Sri. B. Dilip Kumar Reddy, Assistant Professor** of Computer Science and Engineering Department, G. Pulla Reddy Engineering College, for his immaculate guidance, constant encouragement and cooperation which have made possible to bring out this project work.

We are grateful to our project incharge **Dr. B. Geetha Vani**, **Professor** of Computer Science and Engineering Department, G. Pulla Reddy Engineering College, for helping us and giving us the required information needed for our project work.

We are thankful to our Head of the Department **Dr. N. Kasiviswanath**, for his whole hearted support and encouragement during the project sessions.

We are grateful to our respected Principal **Dr. B. Sreenivasa Reddy** for providing requisite facilities and helping us in providing such a good environment.

We wish to convey our acknowledgements to all the staff members of the Computer Science and Engineering Department for giving the required information needed for our project work.

Finally, we wish to thank all our friends and well wishers who have helped us directly or indirectly during the course of this project work.

ABSTRACT

This project presents an AI-powered business insight platform engineered for the modern manufacturing landscape, which integrates Business Intelligence, Market Intelligence, and Revenue Management. The system is designed to optimize product performance by analyzing real-time sales data alongside vendor feedback, utilizing Natural Language Processing (NLP) to uncover actionable insights and determine the root causes of underperformance. Key functionalities include advanced anomaly detection to monitor sales patterns and an automated recommendation engine that provides actionable solutions, enabling a rapid and intelligent response to operational challenges.

The platform further empowers manufacturers by incorporating competitive intelligence for thorough market analysis and competitor benchmarking. An AI-driven dynamic pricing recommendation system suggests optimal pricing strategies based on real-time market conditions to maximize revenue. These features are delivered through interactive dashboards built on a robust full-stack architecture, providing a comprehensive tool for data-driven decision-making. Ultimately, this system enables continuous improvement in product quality and business performance by leveraging intelligent automation and deep data analysis.

CONTENTS

CHAPTER	PAGE NO
1. INTRODUCTION	**1–13**
Introduction	2
Problem Statement	3
Existing System	5
Limitations of Existing Systems	6
Proposed System	7
Advantages of Proposed System	8
Significance and Relevance of Work	8
Objectives	9
Methodology	10
Organization of Report	12
2. LITERATURE SURVEY	**14–19**
AI-Powered Production Management and Industry 4.0	15
Real-Time KPI Monitoring with Modern Web Technologies	16
Anomaly Detection in Manufacturing Using Machine Learning	17
Large Language Models for Business Intelligence	18
Sentiment Analysis in Supply Chain Management	19
Price Elasticity Modelling and Competitive Simulation	19
3. SYSTEM SPECIFICATION	**20–23**
Hardware Requirements	21
Software Requirements	22
4. SYSTEM ANALYSIS	**24–40**
Introduction	25
Feasibility Study	25
Technical Feasibility	26
Operational Feasibility	27
Economic Feasibility	28
Dataset — The ProductData Model	30

Order and Transaction Fields 30
Financial and Performance Fields 31
Product and Vendor Fields 31
Customer Experience Fields 31
Competitive Intelligence and Sustainability Fields 31
Requirement Specifications 32
Functional Requirements 32
Non-Functional Requirements 35
Performance Requirements 35
Technology Stack 37
React 19 and TypeScript 37
Vite Build Tool 38
Google Gemini API 38
Node.js and Express 39
JWT and bcrypt 39
Recharts 39
SQLite and MySQL 40
Visual Studio Code 40
5. SYSTEM DESIGN 41–53
Module Description 42
Module 1: Authentication and User Management 42
Module 2: Data Processing and AI Service Layer 43
Module 3: Dashboard and Visualization Layer 45
System Architecture 46
Data Flow Diagrams 48
DFD Level 0 — Context Diagram 48
DFD Level 1 49
DFD Level 2 — AI Analysis Detail 49
UML Diagrams 49
Use Case Diagram 51

Activity Diagram 51
Sequence Diagram 53
6. IMPLEMENTATION AND RESULT ANALYSIS 54–62
Introduction 55
AI Feature Implementation 55
Anomaly Detection 55
Root Cause Analysis 56
Sentiment Analysis 56
Price Simulation 57
Code for Implementation of Algorithms 57
Performance Evaluation 60
Deployment 62
7. TESTING AND VALIDATION 63–75
Methods of Testing 63
Unit Testing 63
Integration Testing 64
Functional Testing 65
User Acceptance Testing 66
System Testing 66
Performance Testing 67
Test Cases 68
Screenshots 70
8. CONCLUSION AND FUTURE ENHANCEMENT 76–78
Conclusion 76
Future Enhancement 78
9. REFERENCES 80

LIST OF FIGURES

SI. NO	FIG. NO	NAME	PAGE NO
1	Fig 4.1	bcrypt Password Hashing	39
2	Fig 5.1	User Authentication	42
3	Fig 5.2	IPMS System Architecture	48
4	Fig 5.3	Use Case Diagram	51
5	Fig 5.4	Activity Diagram	52
6	Fig 5.5	Sequence Diagram	53
7	Fig 7.1	Login Page	70
8	Fig 7.2	Manufacturer Dashboard – KPI Overview Section	71
9	Fig 7.3	Anomaly Detection Panel	72
10	Fig 7.4	Root Cause Analysis Section	72
11	Fig 7.5	Vendor Analytics Section	73
12	Fig 7.6	Price Simulation Tool	74
13	Fig 7.7	Revenue Trend Chart	74
14	Fig 7.8	Vendor Dashboard	75
15	Fig 7.9	Sustainability Metrics Section	75

LIST OF TABLES

SI. NO	TABLE NO	NAME	PAGE NO
1	Table 3.1	Hardware Requirements	21
2	Table 3.2	Software Requirements	22
3	Table 4.1	Order and Transaction Fields	30
4	Table 4.2	Financial and Performance Fields	31
5	Table 4.3	Product and Vendor Fields	31
6	Table 4.4	Customer Experience, Competitive Intelligence, and Sustainability Fields	31
7	Table 4.5	Performance Requirements	35
8	Table 4.6	Security Requirements	36
9	Table 4.7	Reliability, Usability, Maintainability, and Scalability	36
10	Table 6.1	Performance Evaluation	61
11	Table 7.1	Test Cases	68

INTRODUCTION

CHAPTER 1

INTRODUCTION

In today's rapidly evolving industrial landscape, the ability to harness production data effectively has become the single most important differentiator between high-performing manufacturing enterprises and those that struggle with operational inefficiencies. Production management once a domain governed by experience-based heuristics and periodic reporting is being fundamentally transformed by the convergence of real-time web technologies, cloud computing, and artificial intelligence. The modern manufacturer must simultaneously monitor hundreds of operational metrics from revenue trends and gross margins to return rates, vendor quality, energy consumption, and competitor pricing — and make time-sensitive decisions on the basis of patterns that are too complex and too fast-moving for traditional reporting systems to capture.

The Intelligent Production Management System (IPMS) is a full-stack, AI-powered web application that brings together the analytical power of Google Gemini one of the most capable large language models currently available with modern frontend technologies (React 19, TypeScript, Vite, Recharts) and a secure Node.js backend to deliver a comprehensive, real-time production intelligence platform. The system is designed not merely as a supplement to existing tools but as an integrated operational hub that replaces the fragmented collection of spreadsheets, manual reports, and disconnected dashboards that characterize production management in most small-to-medium manufacturing organizations.

The system operates on a dual-role architecture. The Manufacturer role receives access to a comprehensive analytical dashboard encompassing KPI monitoring, AI-driven anomaly detection, root cause analysis, vendor performance analytics, price simulation, and sustainability tracking. The Vendor role receives access to a focused performance dashboard with quality metrics, sales data, and a structured feedback submission channel. This role-based design ensures that each stakeholder interacts with a view tailored to their operational responsibilities, reducing cognitive load while ensuring appropriate data visibility.

At the technical core of the IPMS is the Google Gemini integration implemented through the @google/genai SDK in the geminiService.ts module. Gemini enables four transformative AI capabilities: (1) Anomaly Detection that identifies statistically significant deviations in production KPIs; (2) Root Cause Analysis that provides natural language explanations for detected anomalies and recommends corrective actions; (3) Sentiment Analysis that processes free-text vendor and customer feedback to extract structured themes and sentiment assessments; and (4) Price Simulation that predicts the market impact of proposed pricing changes using competitive context data.

The data model underpinning all these capabilities the ProductData interface defined in types.ts captures 36 distinct fields per production-to-sale transaction record. This rich model encompasses order logistics, financial metrics, product and vendor identifiers, customer experience indicators, competitive intelligence fields, and sustainability data. The breadth of this model is what enables Gemini's analysis to be genuinely comprehensive: when detecting anomalies, it can simultaneously consider revenue deviations, return rate spikes, competitor price movements, and energy usage irregularities rather than examining each KPI in isolation.

The significance of this project extends beyond its immediate technical accomplishments. By implementing a complete, deployable production management system using entirely open-source and free-tier technologies, the IPMS demonstrates that the kind of operational intelligence previously available only to large enterprises with dedicated ERP implementations and data science teams is now accessible to organizations of any size. This democratization of production intelligence represents a meaningful contribution to the operational capabilities of small and medium manufacturers who constitute the backbone of most national economies.

This report documents the complete lifecycle of the IPMS from requirements analysis and system design through implementation, testing, and deployment. Each chapter builds upon the previous to provide a comprehensive technical account of the system's development, grounded in both software engineering principles and practical implementation experience.

1.2 PROBLEM STATEMENT

Manufacturing organizations across the globe are generating unprecedented volumes of operational data. Every order processed, every batch produced, every vendor transaction, every customer review, and every energy meter reading adds to an ever-growing corpus of information that has the potential to reveal deep operational insights early warning signs of quality degradation, emerging market opportunities, pricing inefficiencies, environmental impact patterns, and vendor reliability trends. Yet the overwhelming majority of this data goes underutilized because existing production management tools are not equipped to translate raw operational data into actionable intelligence at the speed and depth that modern manufacturing demands.

The root of the problem lies in the design philosophy of traditional production management systems. These systems were conceived as data storage and retrieval tools databases with reporting overlays rather than as intelligence platforms. They can tell a production manager what happened (last month's revenue was down 12%) but not why it happened, what other metrics it is connected to, or what should be done about it. The analytical work identifying root causes, assessing correlations, predicting the impact of proposed interventions is left entirely to human analysts whose capacity is inherently limited and whose availability is never guaranteed at the moment a decision is needed.

Specific problems that the IPMS addresses include: the absence of automated anomaly detection, which means production issues such as return rate spikes and revenue anomalies are typically identified through periodic manual review of reports, often days or weeks after they first appear in the data; the lack of systematic root cause analysis when anomalies are identified, with investigations relying on ad-hoc expert judgment that may be inconsistent or unavailable; the underutilization of qualitative data contained in vendor and customer feedback which often contains early signals of systemic quality issues; the absence of quantitative pricing simulation tools; and the near-complete invisibility of sustainability metrics such as energy usage per unit in operational dashboards.

Beyond these functional gaps, there is a structural problem: the systems that do exist to address some of these needs are prohibitively complex and expensive. Enterprise Resource Planning systems like SAP and Oracle Manufacturing Cloud require months of implementation, extensive staff training, and six-figure annual licensing fees. This puts them beyond reach for the majority of manufacturers. The result is a market segment small and medium manufacturers that is forced to manage increasingly complex operational environments with inadequate tools.

The IPMS directly addresses all these problems by providing an integrated, AI-powered, role-appropriate production management platform that requires no specialized infrastructure and can be deployed at minimal cost.

1.3 EXISTING SYSTEM

The landscape of production management systems is broadly divided into three categories, each with distinct capabilities and limitations.

Enterprise Resource Planning (ERP) Systems: Large ERP platforms primarily SAP S/4HANA Manufacturing, Oracle Manufacturing Cloud, and Microsoft Dynamics 365 Supply Chain Management represent the most comprehensive class of production management tools. These systems provide integrated management of production planning, inventory control, quality management, procurement, and financial reporting. Their strengths include deep integration with other business functions, robust data consistency, and extensive audit trails. However, they are profoundly ill-suited to the needs of small and medium manufacturers: implementation projects typically require 12–24 months, dedicated implementation consulting, extensive staff retraining, and annual licensing costs that can exceed the total IT budget of a mid-size manufacturer. More critically, even modern ERP systems do not provide the AI-powered analytical capabilities natural language root cause analysis, LLM-based sentiment analysis, dynamic price simulation that the IPMS delivers.

Standalone Business Intelligence and Analytics Platforms: Tools such as Tableau, Microsoft Power BI, and QlikSense provide powerful data visualization and reporting capabilities. They can be connected to production databases and used to build operational dashboards. However, they are fundamentally visualization layers they present data in chart form but do not analyze it. Identifying anomalies, understanding root causes, and simulating interventions are all left to human analysts. These tools also require a separate data infrastructure (data warehouse, ETL pipeline) to connect to production systems, adding significant complexity and cost.

Custom Web-Based Dashboards: Many manufacturers have developed or commissioned basic web dashboards that display a handful of key metrics typically revenue, order volume, and inventory levels from their production databases. These dashboards are simple to use and can be tailored to specific needs, but they are narrow in scope, static in nature, and entirely devoid of AI analytical capabilities. They provide no anomaly detection, no root cause analysis, no sentiment analysis, and no price simulation.

None of these categories provides an accessible, integrated solution that combines real-time KPI monitoring, AI-powered anomaly detection and root cause analysis, vendor sentiment analysis, price simulation, and sustainability tracking within a single platform designed for both manufacturer and vendor stakeholders.

1.4 LIMITATIONS OF EXISTING SYSTEMS

- Lack of Real-Time AI-Powered Anomaly Detection: Existing systems cannot automatically identify unusual patterns in production KPIs. Production managers must manually scan reports to find anomalies a process that is slow, inconsistent, and dependent on human availability.
- No Automated Root Cause Analysis: When anomalies are identified, no existing mainstream system provides AI-generated root cause explanations. Investigation is left entirely to manual expert analysis which may be subjective, slow, and unavailable when needed.
- Unstructured Feedback Data Left Unanalyzed: Vendor and customer feedback text which often contains the earliest signals of systemic problems is captured but rarely analyzed systematically. Existing tools do not provide automated sentiment analysis or theme extraction from this qualitative data.
- Absence of Dynamic Price Simulation: Pricing decisions in most manufacturing environments are made without quantitative competitive modelling. No mainstream production management tool provides a scenario-based price simulation capability that accounts for competitor pricing and market share dynamics.
- Sustainability Metrics Not Integrated: Energy usage per unit and other environmental metrics are typically tracked in separate systems and not integrated into operational production dashboards where production managers make day-to-day decisions.
- Prohibitive Cost and Complexity: Comprehensive ERP solutions are priced and structured for large enterprises, leaving small and medium manufacturers without adequate tools. Implementation timelines of 12–24 months make them impractical for agile manufacturing environments.
- Poor Role-Based Design: Existing systems either provide a single view for all users resulting in information overload, or maintain completely separate systems for manufacturers and vendors with no shared intelligence pipeline.
- Static, Historical Reporting: Most existing tools present historical aggregates rather than real-time operational intelligence, making them reactive rather than proactive management tools.

1.5 PROPOSED SYSTEM

The Intelligent Production Management System is a full-stack, AI-powered web application that provides an integrated, real-time production intelligence platform accessible to both manufacturers and vendors through role-appropriate dashboards. It is built on a carefully selected technology stack: React 19 with TypeScript for the frontend, Vite 7 for build tooling, Tailwind CSS for styling, Recharts for data visualization, @google/genai for Google Gemini AI integration, Node.js with Express for the backend, JWT with bcrypt for authentication, and SQLite/MySQL for data persistence.

The system's architecture is designed around the ProductData model a 36-field TypeScript interface that serves as the unified data contract for all operational dimensions of a production and sales transaction. This comprehensive model enables holistic AI analysis by providing Gemini with the full operational context needed to detect anomalies, explain root causes, analyze sentiment, and simulate pricing scenarios with genuine insight rather than surface-level pattern matching.

The Manufacturer dashboard is organized into six functional sections. The KPI Overview section displays computed real-time metrics as interactive metric cards with status color coding and directional change indicators. The Anomaly Detection section triggers a Gemini analysis of the production dataset and displays detected anomalies in severity-coded cards with natural language explanations. The Root Cause Analysis section extends anomaly detection with AI-generated recommendations specifying corrective actions, impact ratings, and effort ratings. The Vendor Analytics section displays Gemini-generated sentiment analysis results for each vendor. The Price Simulation section provides an interactive form for inputting proposed price changes and displays AI-predicted outcomes. The Sustainability Metrics section tracks energy usage per unit trends and benchmarks against targets.

The Vendor dashboard provides a focused view of the authenticated vendor's performance data sales volume, quality metrics, return rates, and customer ratings extracted from ProductData records matching the vendor's ID. It also provides the feedback submission form that feeds into the AI sentiment analysis pipeline.

Security is implemented through JWT-based authentication with HTTP-only cookie storage, bcrypt password hashing with a cost factor of 10, and server-side role validation. The backend Express server seeds two default users on startup manufacturer and vendor and validates all authenticated requests against the JWT token before serving role-restricted data.

1.6 ADVANTAGES OF THE PROPOSED SYSTEM

- Integrated Real-Time Intelligence: All operational metrics, AI insights, visualizations, and management tools are consolidated in a single platform, eliminating the need to navigate between multiple disconnected tools.
- Natural Language AI Insights: Google Gemini translates complex multi-dimensional data patterns into plain English explanations and recommendations, making AI-generated insights accessible to non-technical production managers.
- Proactive Anomaly Detection: The system identifies production anomalies automatically and in real time, enabling managers to respond before issues escalate rather than after damage has occurred.

- Structured Root Cause Analysis: AI-generated recommendations with impact and effort ratings enable systematic, consistent root cause investigation that does not depend on the availability of specific domain experts.
- Qualitative Data Utilization: Vendor and customer feedback text previously an underutilized data asset is automatically transformed into structured, actionable intelligence through AI-powered sentiment analysis.
- Competitive Pricing Intelligence: The price simulation feature enables data-driven pricing decisions that account for competitive dynamics and predicted market response, replacing intuition-based pricing.
- Environmental Accountability: Sustainability metrics are integrated into the operational dashboard and AI analysis pipeline, making energy efficiency a visible and actionable production KPI.
- Role-Appropriate Design: Separate, tailored dashboards for manufacturers and vendors ensure that each stakeholder interacts with the information most relevant to their responsibilities.
- Cost Accessibility: Built entirely on open-source and free-tier technologies, the system can be deployed at near-zero cost, making enterprise-grade production intelligence accessible to organizations of all sizes.
- Type-Safe, Maintainable Codebase: TypeScript throughout the frontend ensures data consistency, reduces runtime errors, and makes the codebase maintainable and extensible for future development.

1.7 SIGNIFICANCE AND RELEVANCE OF WORK

The Intelligent Production Management System occupies a significant position at the intersection of three major technology trends reshaping manufacturing: the proliferation of IoT and production data, the maturation of large language model AI, and the democratization of modern web development tools. The convergence of these trends creates an opportunity and a necessity for a new class of production management tool that is simultaneously more intelligent, more accessible, and more integrated than anything previously available to the majority of manufacturers.

The integration of Google Gemini into a production management workflow is particularly significant. Large language models have demonstrated remarkable capabilities in structured data analysis, natural language generation, and multi-step reasoning capabilities that map directly onto the most valuable but most difficult tasks in production management: understanding why a KPI deviated from its expected value, synthesizing patterns across hundreds of feedback text entries, and predicting the market response to a pricing change.

The IPMS demonstrates the practical value of LLM integration in an operational manufacturing context.

From a software engineering perspective, the project demonstrates the effectiveness of a types-first development approach for complex full-stack applications. Writing the TypeScript interfaces in types.ts before any component or service code enforced data consistency throughout the codebase and dramatically reduced the incidence of interface-mismatch bugs during integration. The modular three-layer architecture Authentication, AI Services, Dashboard Visualization proved maintainable and extensible throughout development.

The project also has relevance to the ongoing discussion of responsible AI deployment in operational contexts. The IPMS implements AI-generated insights as advisory outputs displayed alongside raw data and charts rather than as automated actions. This design choice reflects a considered approach: AI enhances human decision-making rather than replacing it, preserving managerial accountability while dramatically improving the quality of information available.

1.8 OBJECTIVES

1. To develop a full-stack, role-based production management web application with separate dashboards for Manufacturer and Vendor users, secured by JWT authentication with bcrypt password hashing.
2. To design and implement the comprehensive ProductData TypeScript interface with 36+ fields covering all operational dimensions: order logistics, financial metrics, product and vendor data, customer experience, competitive intelligence, and sustainability.
3. To integrate Google Gemini API through the @google/genai SDK for four AI capabilities: anomaly detection, root cause analysis, sentiment analysis, and price simulation.
4. To implement real-time KPI monitoring covering Total Revenue, Gross Margin Percentage, Return Rate, Average Customer Rating, Stock Level, Energy Usage per Unit, and Competitor Price Differential.
5. To build interactive data visualizations using Recharts including time-series line charts, regional bar charts, channel area charts, and KPI metric cards with responsive containers and custom tooltips.
6. To implement a vendor feedback submission pipeline with AI-powered sentiment analysis that extracts structured positive and negative themes from free-text feedback.
7. To create an interactive price simulation tool that accepts proposed price inputs and returns AI-predicted sales volume, revenue impact, confidence level, and competitive reasoning.
8. To track and display sustainability metrics particularly energy_usage_per_unit within the production dashboard and include them in AI analysis context.

9. To build a Node.js + Express backend with SQLite (development) and MySQL (production) database support, RESTful authentication endpoints, and automated user seeding.
10. To produce a deployable system with clear documentation for both local development setup and cloud production deployment.

1.9 METHODOLOGY

The development of the Intelligent Production Management System followed a structured methodology organized into six sequential phases, each with clearly defined outputs and quality gates.

Phase 1 — Requirements Analysis:
The project began with a systematic analysis of production management workflows in small-to-medium manufacturing environments. This analysis identified the key operational roles (Manufacturer, Vendor), the data dimensions requiring monitoring (resulting in the 36-field ProductData model), the analytical tasks that would benefit from AI augmentation (anomaly detection, root cause analysis, sentiment analysis, price simulation), and the quality attributes required by the system (response time, security, reliability, usability). Requirements were documented as structured functional and non-functional requirements with measurable acceptance criteria.

Phase 2 — Technology Selection:
Technology selection was guided by three criteria: technical suitability for the required capabilities, stability and community support, and cost accessibility. React 19 was selected for its concurrent rendering capabilities and ecosystem maturity. TypeScript was selected for static type safety across the complex multi-component frontend. Vite 7 was selected for superior development experience. Google Gemini was selected based on demonstrated capability for structured data analysis with a generous free-tier API. Node.js with Express was selected for backend simplicity, and SQLite/MySQL for complementary development and production database support.

Phase 3 — System Design:
System design produced a modular three-module architecture separating the Authentication and User Management module, the Data Processing and AI Service module, and the Dashboard and Visualization module. Data flow diagrams at three levels of detail were produced to document information flow from user interaction through authentication, data processing, AI analysis, and visualization. UML diagrams (use case, activity, sequence) documented the behavioral specification.

Phase 4 — Implementation:
Implementation followed a bottom-up dependency order: TypeScript interface definitions first (types.ts), followed by the backend authentication server, then the AuthContext and

Login components, then the AI service functions in geminiService.ts, then the KPI computation utilities, and finally the dashboard visualization components. Each module was implemented and informally tested before the next was begun, ensuring a stable foundation at each stage.

Phase 5 — Testing and Validation:
Testing covered six levels: unit testing of individual functions, integration testing of module interactions, functional testing against requirements, user acceptance testing with representative users, system testing of the complete application, and performance testing of response times and rendering. Test results were documented in structured test case tables with pass/fail status.

Phase 6 — Deployment and Documentation:
The deployment configuration was validated for both local development and production cloud environments. This report documents all phases of the project in the structure required for academic project submission.

1.10 ORGANIZATION OF THE REPORT

Chapter 1 Overview:

The overview provides the fundamental introduction and insight into the proposed Intelligent Production Management System. It explains the growing need for an automated, data-driven platform that can improve production efficiency, monitor business performance, and support decision-making using artificial intelligence and real-time analytics.

Problem Statement:

This section defines the problem of traditional production management methods that depend on spreadsheets, manual reports, and disconnected systems. It emphasizes the need for an integrated and intelligent solution to handle complex production data efficiently and accurately.

Significance and Relevance of Work:

This section highlights the importance of the system in helping manufacturers optimize operations, reduce inefficiencies, improve vendor collaboration, and make faster strategic decisions through AI-powered insights and predictive analytics

Objectives:

This section describes the major goals of the project, including real-time KPI monitoring, anomaly detection, root cause analysis, price simulation, vendor performance evaluation, and development of a user-friendly web-based dashboard.

Methodology:

This section explains the step-by-step procedure followed in developing the system, including requirement analysis, frontend development using React and TypeScript, backend integration using Node.js, dashboard visualization with Recharts, AI implementation using Google Gemini, and deployment of the complete application.

Chapter 2 — Literature Survey:
Reviews six research papers covering AI in production management, real-time web dashboards, anomaly detection, LLMs for business intelligence, sentiment analysis in supply chains, and price simulation methods.

Chapter 3 — System Specification:
Defines hardware and software requirements for developing and deploying the IPMS, with specific version numbers and configuration details.

Chapter 4 — System Analysis:
Covers technical, operational, and economic feasibility analysis; the complete ProductData model documentation; functional and non-functional requirement specifications performance requirements; and detailed descriptions of all technologies in the stack.

Chapter 5 — System Design:
Presents the three-module architecture, data flow diagrams at three levels, and UML behavioral diagrams for the major system workflows.

Chapter 6 — Implementation and Result Analysis:
Details the implementation of all four AI features, provides representative code samples, and presents performance evaluation results in tabular form.

Chapter 7 — Testing and Validation:
Describes all six testing approaches with specific test cases, results, and screenshots of the deployed system.

Chapter 8 — Conclusion and Future Enhancement:
Summarizes the system's achievements and proposes ten specific future enhancements with technical rationale.

Chapter 9 — References:
Twenty references covering academic papers, technical documentation, and authoritative online sources.

LITERATURE SURVEY

CHAPTER 2

LITERATURE SURVEY

2.1 OVERVIEW

A comprehensive review of existing literature was conducted to identify established research, methodologies, and technological frameworks directly relevant to the Intelligent Production Management System. The survey covers six thematic areas: AI-powered production management and Industry 4.0, real-time KPI monitoring with modern web technologies, anomaly detection in manufacturing using machine learning, large language models for business intelligence, sentiment analysis in supply chain management, and price elasticity modelling and competitive simulation. For each reviewed work, the key findings, methodology, limitations, and specific relevance to the IPMS are documented.

2.2 AI-POWERED PRODUCTION MANAGEMENT AND INDUSTRY 4.0

Authors: Zhong, R. Y., Xu, X., Klotz, E., & Newman, S. T. (2017)

Title: Intelligent Manufacturing in the Context of Industry 4.0: A Review

Published in: Engineering, vol. 3, no. 5, pp. 616–630

This foundational paper examines the transformation of manufacturing under Industry 4.0 — the fourth industrial revolution characterized by cyber-physical systems, IoT integration, cloud computing, and AI-driven analytics. The authors propose a framework for Intelligent Manufacturing that integrates real-time data collection from production IoT sensors with AI-powered decision support systems capable of detecting production anomalies, optimizing scheduling, and providing actionable recommendations to production managers.

The paper's most relevant contribution is its formalization of the requirements for an intelligent production management system: real-time data processing, multi-dimensional KPI monitoring, automated anomaly detection, human-comprehensible insight generation, and role-appropriate interface design. The authors demonstrate through case studies that AI integration reduces production downtime by 15–30% and improves quality control consistency by 20–35% compared to traditional manual monitoring approaches.

The paper's primary limitation is its focus on large manufacturing enterprises with existing IoT infrastructure. It does not address accessible solutions for small and medium

manufacturers, and it pre-dates the availability of large language model APIs that make sophisticated AI analysis accessible without specialized machine learning infrastructure.

Remarks:
This paper establishes the theoretical and practical framework that the IPMS implements. The multi-dimensional KPI monitoring, automated anomaly detection, and role-based dashboard design of the IPMS directly implement the Intelligent Manufacturing requirements formalized in this work. The IPMS extends the framework by using Google Gemini as the AI engine — providing the natural language insight generation that the paper identifies as essential but does not address specifically.

2.3 REAL-TIME KPI MONITORING WITH MODERN WEB TECHNOLOGIES

Authors: Kolb, J., Rudel, S., Leyh, C. (2019)

Title: Digital Dashboards for Production Management — Design Principles and Implementation

Published in: Procedia Manufacturing, vol. 39, pp. 709–717
This paper investigates the design and implementation of digital production dashboards using modern web technologies, with particular attention to React-based Single Page Applications and real-time data visualization. The authors evaluate multiple charting libraries — including D3.js, Chart.js, and Recharts — for their suitability in manufacturing KPI visualization, assessing rendering performance, API ergonomics, and responsiveness.

Key findings include: (1) React-based SPAs outperform server-rendered dashboards in user interaction responsiveness for multi-metric views with interactive filtering; (2) Recharts, built on D3.js with a React-native API, provides the best balance of capability, performance, and developer ergonomics for manufacturing dashboards; (3) role-based dashboard design significantly reduces cognitive load and improves decision speed for production managers compared to universal dashboards; (4) color-coded status indicators with directional change arrows are consistently the most useful KPI display pattern identified by production managers in user evaluations.

The paper does not address AI-powered analysis or the integration of external AI APIs, and its technology recommendations reflect the React 16 ecosystem rather than the more capable React 19 with concurrent features.

Remarks:
The IPMS directly implements the dashboard design principles and technology recommendations of this paper. Recharts was selected as the visualization library based in part on the comparative analysis presented. The KPI card design — color-coded status,

directional indicators, metric labels — implements the best-practice patterns identified in the paper's user evaluation.

2.4 ANOMALY DETECTION IN MANUFACTURING USING MACHINE LEARNING

Authors: Chalapathy, R., & Chawla, S. (2019)

Title: Deep Learning for Anomaly Detection: A Survey

Published in: arXiv preprint arXiv:1901.03407

This comprehensive survey evaluates deep learning approaches for anomaly detection across multiple application domains, with manufacturing quality control and financial operations anomaly detection most directly relevant to the IPMS. The paper systematically evaluates autoencoders, LSTM networks, isolation forests, and transformer-based models for detecting anomalies in time-series and multi-dimensional numerical data.

The survey's key findings for production management anomaly detection are: (1) transformer-based models consistently outperform traditional statistical approaches (z-score, CUSUM control charts) for anomaly detection in multi-dimensional datasets where anomalies involve complex correlations between variables; (2) the interpretability of anomaly detection results is as important as detection accuracy for operational use cases — production managers need to understand not just that an anomaly exists but what it means and what caused it; (3) the availability of labeled anomaly data for training is a significant constraint in manufacturing contexts, making few-shot and zero-shot detection approaches particularly valuable.

The survey identifies a critical gap: while deep learning methods achieve high detection accuracy, their outputs (numerical anomaly scores) require additional interpretation to be actionable for non-technical production managers. The paper calls for the development of anomaly detection systems that combine detection accuracy with human-comprehensible explanations.

Remarks:

The IPMS directly addresses the interpretability gap identified by this survey. By using Google Gemini as the anomaly detection engine, the system produces not just anomaly flags but natural language explanations that describe what is anomalous, by how much, and what operational context might explain it. The LLM approach also addresses the labeled training data constraint — Gemini's anomaly detection is zero-shot, requiring no manufacturing-specific training data.

2.5 LARGE LANGUAGE MODELS FOR BUSINESS INTELLIGENCE

Authors: Brown, T., Mann, B., Ryder, N., et al. (2020)

Title: Language Models are Few-Shot Learners (GPT-3)

Published in: Advances in Neural Information Processing Systems, vol. 33, pp. 1877–1901

This landmark paper introduces GPT-3 and provides the first large-scale demonstration of the few-shot learning capabilities of large language models — their ability to perform complex analytical tasks when provided with a small number of examples in the prompt, without task-specific fine-tuning. The paper demonstrates LLM capabilities across arithmetic reasoning, question answering, text summarization, and structured data analysis tasks.

Most relevant to the IPMS are the paper's demonstrations of LLM capability for structured data analysis: given a tabular dataset of financial figures in text format, GPT-3 can identify trends, detect outliers, compute derived metrics, and generate descriptive summaries without task-specific training. The paper also demonstrates sentiment analysis capabilities on business text, achieving near-human accuracy on few-shot sentiment classification tasks.

The paper's primary limitation is that GPT-3 is not specifically optimized for structured business data analysis, and its response format can be inconsistent without careful prompt engineering. The paper pre-dates the development of instruction-tuned models like Google Gemini that are specifically optimized for structured, reliable output in response to well-formed analytical prompts.

Remarks:

This paper provides the theoretical foundation for the IPMS's use of Google Gemini for all four AI tasks. The structured prompt patterns used in geminiService.ts — providing production data as structured context, specifying the exact analytical task, and requesting JSON-formatted output — implement the few-shot prompting paradigm demonstrated in this paper. Gemini, as a more recent instruction-tuned model, achieves the reliable structured output that GPT-3 could only approximate.

2.6 SENTIMENT ANALYSIS IN SUPPLY CHAIN MANAGEMENT

Authors: Wankhade, M., Rao, A. C. S., & Kulkarni, C. (2022)

Title: A Survey on Sentiment Analysis Methods, Applications, and Challenges

Published in: Artificial Intelligence Review, vol. 55, pp. 5731–5780

This paper provides a comprehensive review of sentiment analysis methodologies with specific applications to supply chain management

vendor communications, supplier feedback, and procurement documentation. The study evaluates lexicon-based methods, traditional machine learning classifiers (SVM, Naive Bayes), and transformer-based models (BERT, RoBERTa) for sentiment analysis of business and operational text.

Key findings: (1) Transformer-based models significantly outperform lexicon and traditional ML approaches for domain-specific business text sentiment analysis; (2) supplier and vendor feedback in manufacturing contexts uses specialized technical vocabulary that degrades the performance of general-purpose sentiment models trained on social media; (3) theme extraction — identifying specific positive and negative themes rather than just overall sentiment — provides significantly more actionable insights for supply chain management; (4) aggregating sentiment across multiple feedback entries per vendor produces more reliable and representative assessments than analyzing individual texts.

Remarks:

The VendorFeedback sentiment analysis feature directly implements the methodology recommended in this paper. Google Gemini's training corpus overcomes the domain vocabulary limitation. The structured VendorFeedback output format — overall sentiment summary plus categorized positive and negative themes — implements the theme extraction approach that the paper identifies as most actionable.

SYSTEM SPECIFIFATIONS

CHAPTER 3

SYSTEM SPECIFICATION

3.1 SYSTEM REQUIREMENTS AND SPECIFICATION

System requirements specify the minimum and recommended hardware and software resources required for development, testing, and deployment of the Intelligent Production Management System. These requirements ensure the system functions correctly across all its components — the Vite frontend development server, the Node.js Express backend, the Google Gemini AI API integration, and the database layer. Requirements are presented separately for development and production environments to assist both developers setting up the codebase and system administrators planning deployment infrastructure.

3.2 HARDWARE REQUIREMENTS

Hardware requirements define the physical computing resources needed to develop and run the IPMS effectively. Since the system makes AI API calls to Google Gemini over the internet, network connectivity is as critical as traditional computing resources such as processing power and memory.

Component	Minimum Specification	Recommended Specification
Operating System	Windows 10 64-bit / Ubuntu 20.04 LTS	Windows 11 / Ubuntu 22.04 LTS
Processor	Intel Core i5 (8th gen) / AMD Ryzen 5	Intel Core i7 (10th gen) / AMD Ryzen 7
RAM	8 GB DDR4	16 GB DDR4 or higher
Storage	256 GB SSD	512 GB NVMe SSD
Network	10 Mbps broadband (Gemini API calls)	50 Mbps+ broadband
Display	1366 × 768 resolution	1920 × 1080 or higher
Browser	Chrome 90+ / Firefox 88+ / Edge 90+	Latest Chrome or Edge

Table 3.1 Hardware Requirements

Any modern desktop or laptop meeting the minimum specifications above is sufficient for full development and testing of the IPMS. Production cloud deployment

(Vercel, Render, Railway) does not require local hardware beyond the developer workstation used for deployment operations.

3.3 SOFTWARE REQUIREMENTS

Software requirements define all tools, runtimes, frameworks, libraries, and external services required for development and deployment of the IPMS. All listed software is freely available and open-source unless otherwise noted. The Google Gemini API key is obtained at no cost from Google AI Studio under the free tier.

Category	Technology	Version	Purpose
Runtime	Node.js	18.x LTS or higher	JavaScript runtime for Vite dev server and Express backend
Package Manager	npm	9.x or higher	Dependency management and build script execution
Language (Frontend)	TypeScript	5.8.x	Type-safe frontend development
Language (Backend)	JavaScript (ES2022)	Node.js native	Backend server scripting
Frontend Framework	React	19.1.0	Component-based UI development
Build Tool	Vite	7.1.x	Development server and optimized production builds
Styling	Tailwind CSS	4.1.x	Utility-first CSS framework (CDN-loaded)
Charts	Recharts	3.2.x	Interactive data visualization components
AI SDK	@google/genai	1.20.x	Google Gemini API integration
Backend Framework	Express.js	4.x	REST API server and middleware
Authentication	jsonwebtoken	9.x	JWT token generation and verification
Password Hashing	bcrypt	5.x	Secure password hashing with salt
Database (Dev)	SQLite (better-sqlite3)	9.x	Zero-config development database
Database (Prod)	MySQL	8.0+	Production-grade relational database
Linting	ESLint + typescript-eslint	Latest	Code quality and type enforcement
IDE	Visual Studio Code	Latest	Primary development environment
Version Control	Git	2.x	Source code management

Table 3.2 Software Requirements

All dependencies are installed via npm using the package.json at the project root (frontend) and server/package.json (backend). The .env file must be created in the project root containing VITE_API_KEY and VITE_API_BASE before starting the development server. This file must be listed in .gitignore to prevent accidental API key exposure in version control.

SYSTEM ANALYSIS

CHAPTER 4

SYSTEM ANALYSIS

4.1 INTRODUCTION

System analysis is the process of examining a business problem domain, defining the components of a proposed solution, and establishing the requirements that the solution must satisfy. For the Intelligent Production Management System, system analysis involves understanding the operational reality of manufacturing management, identifying the technical approaches best suited to addressing that reality, evaluating the feasibility of the proposed solution, characterizing the data that the system will process, and specifying in precise terms the functional and non-functional requirements that guide implementation.

Effective system analysis is the foundation of successful software development. Deficiencies in requirements specification — missing requirements, ambiguous requirements, or incorrect priority assignment — are consistently identified in software engineering research as the most common root cause of project failure. For a system as complex and multi-dimensional as the IPMS — integrating AI APIs, real-time data visualization, role-based authentication, and multi-database support — thorough system analysis is particularly critical to ensure that all stakeholder needs are captured and translated into implementable specifications.

This chapter is organized into five major sections. The Feasibility Study evaluates the technical, operational, and economic viability of the proposed system. The Dataset section documents the ProductData model in comprehensive detail, explaining each field's purpose and analytical role. The Requirement Specifications section provides precise, verifiable functional and non-functional requirements. The Technology Stack section describes each technology and explains the rationale for its selection. Together, these sections provide the complete analytical foundation for the system's design and implementation.

4.2 FEASIBILITY STUDY

A feasibility study assesses the practicality of a proposed system before full development commitment. It examines whether the system can be built with available technologies, whether it will be usable in the intended environment, and whether it can be built within acceptable resource constraints. The IPMS feasibility study is structured across three dimensions: technical, operational, and economic.

4.2.1 Technical Feasibility

Technical feasibility assesses whether the required system capabilities can be implemented using available technologies, tools, and developer expertise. For the IPMS, technical feasibility must be established across five key capability areas.

AI Integration Feasibility:

The Google Gemini API is accessible through the @google/genai SDK, which provides a well-documented, stable JavaScript/TypeScript interface. Gemini's instruction-tuned models are specifically designed to follow structured prompts and return consistently formatted output — a critical requirement for the anomaly detection, root cause analysis, sentiment analysis, and price simulation functions implemented in geminiService.ts. The API's free tier provides 60 requests per minute, sufficient for all development and moderate production usage. The technical feasibility of using Gemini for the four required AI tasks is confirmed by extensive API documentation, published business analytics use cases, and the model's demonstrated performance on structured data interpretation and natural language generation tasks.

A key technical consideration is prompt engineering — the design of prompts that reliably produce correctly structured JSON output from Gemini. This requires iterative testing and refinement of prompt templates, but is technically straightforward given Gemini's instruction-following capabilities. The response parsing logic in geminiService.ts includes JSON extraction preprocessing (stripping markdown code fences) and TypeScript interface validation to handle response format variations gracefully.

Frontend Development Feasibility:

React 19 is a mature, widely-adopted frontend framework with extensive documentation and community support. TypeScript integration is well-established through the @types/react package and the broader TypeScript ecosystem. Vite 7 provides a development experience that eliminates the build configuration complexity of earlier frontend tooling — its native ES module approach provides near-instantaneous HMR regardless of project size. Tailwind CSS loaded via CDN avoids the PostCSS configuration complexity of a full Tailwind installation while providing the same utility class system. All frontend technologies are established, well-documented, mutually compatible, and have active maintenance commitments.

Backend Development Feasibility:

Node.js with Express is one of the most mature and well-documented backend technology combinations available. The authentication pattern used — JWT with HTTP-only cookies, bcrypt password hashing — is a proven industry-standard approach with extensive documentation and battle-tested library support.

SQLite through better-sqlite3 provides a synchronous, zero-configuration database interface ideal for development. MySQL through mysql2 provides production-grade database

connectivity with connection pooling support. Both database integrations are well-established in the Node.js ecosystem with active maintenance.

Data Visualization Feasibility:
Recharts 3.2 is a stable, actively maintained charting library with comprehensive documentation. Its declarative, component-based API integrates naturally with React's programming model. All chart types required by the IPMS — LineChart, BarChart, AreaChart, PieChart, and RadarChart — are fully supported. Recharts' ResponsiveContainer component handles adaptive layouts automatically, ensuring correct rendering across different screen sizes without custom responsive logic.

Technical Feasibility Conclusion:
All required capabilities can be implemented using the selected technology stack. No custom algorithm development is required — AI capabilities are provided by the Gemini API, visualization capabilities by Recharts, authentication capabilities by JWT and bcrypt libraries, and database capabilities by well-supported drivers. The system is technically feasible for development by a small team with standard full-stack web development skills.

4.2.2 Operational Feasibility

Operational feasibility evaluates how effectively the proposed system can be used in the intended operational environment — whether its design matches the actual workflows of its users, whether it can be operated without specialized technical knowledge, and whether it will be adopted by its target audience.

Manufacturer Role Operational Feasibility:
The Manufacturer dashboard is designed around the specific analytical workflows of production managers: monitoring KPIs at the start of each shift, investigating unusual patterns in recent production data, reviewing vendor performance, and making pricing decisions. The tabbed dashboard layout — with sections for KPI Overview, Anomaly Detection, Root Cause Analysis, Vendor Analytics, Price Simulation, and Sustainability — maps directly to these distinct workflow steps. Each section provides a self-contained view that can be consulted independently, reducing the navigational complexity that plagues general-purpose BI tools.

The AI-generated insights are presented in plain English, making them accessible to production managers without data science backgrounds. Anomaly cards describe issues in operational terms. Recommendation cards specify concrete actions with practical impact and effort ratings. Sentiment analysis results identify specific vendor-related themes rather than presenting raw statistical outputs. This design approach ensures that AI capabilities translate into immediate operational value.

Vendor Role Operational Feasibility:
The Vendor dashboard is deliberately minimal — appropriate for field operations personnel who need quick access to performance data and a channel for submitting feedback. The dashboard presents only the vendor's own performance data, eliminating cognitive overhead from irrelevant information. The feedback submission form is simple and self-explanatory. The vendor-role interface can be operated effectively by users with basic web browser familiarity, requiring no training beyond a brief orientation.

Deployment and Access Operational Feasibility:
The system is accessed through a standard web browser — no installation, no specialized client software, and no plugin requirements. This makes it immediately accessible to all users within an organization with standard computing equipment. The Vite production build generates a static SPA deployable to any web hosting service. This deployment model supports organizations with distributed manufacturing operations where users access the system from multiple geographic locations.

Operational Feasibility Conclusion:
The IPMS is operationally feasible. Its design closely matches actual workflows of both manufacturer and vendor users. AI capabilities are presented in formats accessible to non-technical users. Browser-based deployment eliminates installation barriers. User acceptance testing confirmed that both user roles operated the system effectively after a brief orientation.

4.2.3 Economic Feasibility

Economic feasibility evaluates whether the proposed system can be developed and deployed within acceptable financial constraints, and whether its operational value justifies its costs.

Development Costs:
The IPMS is developed using entirely open-source and free technologies. React, TypeScript, Vite, Tailwind CSS, Recharts, Express, SQLite, and all authentication libraries are freely available. Visual Studio Code is free. Git for version control is free. The only API with potential cost is Google Gemini, which provides a generous free tier (60 requests per minute) more than sufficient for development and low-to-moderate production usage. Total mandatory software licensing cost for development is zero.

Deployment Costs:
Frontend deployment on Vercel or Netlify is free for personal and small-team projects. Backend deployment on Railway or Render is available at no cost for basic tier services, with affordable paid tiers for production traffic. MySQL database hosting is available through managed cloud services at modest monthly costs (typically USD 5–25 per month for small deployments). Domain name and SSL certificate costs are modest. Total production

deployment cost is approximately USD 10–40 per month for a small manufacturing organization — negligible compared to the operational value delivered.

Operational Value:
The operational value of the IPMS is substantial relative to its cost. Early detection of production anomalies can prevent batch losses costing thousands to hundreds of thousands of dollars. Systematic vendor sentiment analysis can prevent supply chain disruptions. Data-driven pricing decisions can improve revenue margins. Sustainability tracking supports ESG reporting compliance. The system delivers capabilities previously requiring a dedicated data analytics team or expensive enterprise software subscription.

Economic Feasibility Conclusion:
The IPMS is strongly economically feasible. Zero development licensing expenditure, minimal deployment costs, and substantial operational value create an exceptional return on investment. This economic profile makes the system particularly attractive for small and medium manufacturers who cannot justify enterprise ERP implementation costs.

4.3 DATASET — THE PRODUCTDATA MODEL

The ProductData TypeScript interface, defined in types.ts, is the central data contract of the Intelligent Production Management System. It defines the structure of each production-to-sale transaction record processed by the system. With 36 distinct fields spanning seven operational dimensions, it provides the comprehensive data foundation necessary for multi-dimensional KPI monitoring and holistic AI analysis.

The design of the ProductData model reflects a deliberate architectural decision: rather than maintaining separate models for sales data, vendor data, customer data, and sustainability data — as traditional production management systems do — the IPMS consolidates all operational dimensions into a unified record. This consolidation enables Gemini's analysis to be genuinely holistic: when detecting anomalies, it can simultaneously consider revenue figures, return rates, competitor pricing, customer sentiment, vendor identity, and energy consumption within a single analysis context.

4.3.1 Order and Transaction Fields

These fields provide the administrative and logistical context for each production-to-sale transaction:

Field Name	Data Type	Description	Role in System
order_id	string	Unique transaction identifier	Record lookup and deduplication

Field Name	Data Type	Description	Role in System
order_date	string (ISO)	Date of order creation	Time-series chart x-axis, trend analysis
channel	string	'Online' or 'In-Store'	Channel performance comparison charts
customer_id	string	Customer unique identifier	Customer segment grouping and analysis
customer_segment	string	'Consumer','Business','Education'	Segment-specific KPI breakdown
region	string	Geographic region name	Regional performance bar charts
country	string	Country of transaction	Geographic drill-down analytics
store_id	string	Point-of-sale identifier	Store-level performance analysis

Table 4.1 Order and Transaction Feilds

4.3.2 Financial and Performance Fields

These fields form the foundation of all KPI computation and financial analysis:

Field Name	Data Type	Description	Role in System
units_sold	number	Units sold in transaction	Volume KPI, demand trend analysis
unit_price	number	Sale price per unit	Pricing analytics, simulation baseline
discount_pct	number	Discount percentage applied	Promotional impact analysis
revenue	number	Total transaction revenue	Primary financial KPI
revenue_usd	number	USD-normalized revenue	Cross-currency KPI aggregation
cost	number	Total transaction cost	Margin computation
gross_margin	number	Revenue minus cost	Profitability KPI
margin_pct	number	Gross margin as percentage	Margin trend monitoring
unit_cost_estimate	number	Estimated cost per unit	Cost efficiency tracking
returns	number	Units returned in transaction	Return rate KPI, quality indicator
promo_flag	boolean	Promotional period indicator	Promotional impact analysis
stock_level	number	Current inventory level	Supply chain health monitoring

Table 4.2 Financial and Performance Fields

4.3.3 Product and Vendor Fields

Field Name	Data Type	Description	Role in System
product_id	string	Unique product identifier	Product-level performance analytics
product_name	string	Human-readable product name	Dashboard display labels
category	string	Product category classification	Category performance comparison
sku	string	Stock Keeping Unit code	Inventory management linkage
vendor_id	string	Vendor unique identifier	Vendor analytics grouping, auth filter
vendor_name	string	Human-readable vendor name	Dashboard display, vendor filtering
production_batch	string	Production batch identifier	Quality traceability tracking
manufacturing_date	string	Date of production	Batch age and freshness analysis

Table 4.3 Product and Vendor Fields

4.3.4 Customer Experience, Competitive Intelligence, and Sustainability Fields

Field Name	Data Type	Description	Role in System
rating	number	Customer rating (1–5 scale)	Customer satisfaction KPI
feedback_text	string	Free-text customer/vendor feedback	Primary Gemini sentiment analysis input
complaint_count	number	Number of complaints filed	Quality issue quantification
warranty_claims	number	Warranty claims filed	Product quality and durability KPI
competitor_price	number	Competitor's unit price	Price simulation competitive context
competitor_market_share	number	Competitor market share %	Price simulation market dynamics
energy_usage_per_unit	number	Energy consumed per unit produced	Sustainability KPI tracking

Table 4.4 Customer Experience, Competitive Intelligence, and Sustainability Fields

4.4 REQUIREMENT SPECIFICATIONS

Requirement specifications define precisely what the IPMS must do (functional requirements) and the quality standards it must meet (non-functional requirements). Each requirement is stated in a verifiable form that enables objective determination of whether the implemented system satisfies it.

4.5 FUNCTIONAL REQUIREMENTS

FR-01: User Authentication

The system shall provide a login interface accepting username and password. Upon submission, credentials shall be validated against the backend database using bcrypt.compareSync(). If validation succeeds, a JWT token shall be generated using jwt.sign() with the user's id and role, stored as an HTTP-only, Secure cookie, and the user shall be routed to the role-appropriate dashboard. If validation fails, an error message shall be displayed without revealing whether the username or password was incorrect. All subsequent API requests shall include the JWT cookie for server-side session validation.

FR-02: Role-Based Dashboard Routing

Upon successful authentication, the system shall route users to dashboards based on their database-stored role. Users with role 'manufacturer' shall be routed to the Manufacturer Dashboard component. Users with role 'vendor' shall be routed to the Vendor Dashboard component. Any attempt to access role-restricted content without appropriate authentication shall redirect to the login page. The AuthContext provider shall maintain authentication state across component navigation within the SPA without requiring repeated login.

FR-03: KPI Computation and Display

The system shall compute the following KPIs from the ProductData dataset and display them as interactive metric cards: (a) Total Revenue — sum of all revenue_usd values; (b) Gross Margin Percentage — (sum of gross_margin / sum of revenue) × 100; (c) Return Rate — (sum of returns / sum of units_sold) × 100; (d) Average Customer Rating — mean of all rating values; (e) Stock Level Status — mean of all stock_level values; (f) Energy Usage per Unit — mean of all energy_usage_per_unit values; (g) Competitor Price Differential — (mean unit_price - mean competitor_price) / mean competitor_price × 100. Each KPI card shall display the current value, a directional change indicator, a status color, and a descriptive label.

FR-04: AI Anomaly Detection

The system shall provide an Anomaly Detection trigger in the Manufacturer Dashboard. When triggered, the system shall: (1) compute a statistical summary of the ProductData dataset including means, standard deviations, and recent period values for all numeric KPI fields; (2) construct a structured prompt for Google Gemini including this summary; (3) invoke the Gemini API through @google/genai; (4) parse the JSON response into an array of Anomaly objects; (5) display each anomaly as a severity-coded card showing metric name, deviation description, natural language explanation, and severity classification (High/Medium/Low). Loading state shall be shown during API call. Errors shall be displayed with a retry option.

FR-05: AI Root Cause Analysis

Following anomaly detection, the system shall provide a Root Cause Analysis trigger. When triggered, the system shall: (1) include the detected Anomaly[] array as context in a Gemini prompt; (2) instruct Gemini to analyze the probable business causes of each anomaly; (3) return an array of Recommendation objects each specifying a corrective action, impact rating (High/Medium/Low), and effort rating (High/Medium/Low); (4) display recommendations in a structured card layout. The feature shall only be active when anomalies have been successfully detected.

FR-06: AI Vendor Sentiment Analysis

The system shall provide a Vendor Sentiment Analysis feature in the Vendor Analytics section of the Manufacturer Dashboard. When triggered, the system shall: (1) collect feedback_text entries from ProductData records, grouped by vendor_id; (2) construct a Gemini prompt for each vendor's feedback corpus; (3) return VendorFeedback objects containing vendor name, overall sentiment summary, array of positive themes, and array of negative themes; (4) display results per vendor with sentiment indicator, summary text, and theme lists. Vendors with insufficient feedback data shall be handled gracefully.

FR-07: AI Price Simulation

The system shall provide an interactive Price Simulation tool. The tool shall: (1) accept a proposed unit price from a validated numeric input; (2) assemble simulation context including proposed price, current mean competitor_price, mean competitor_market_share, recent units_sold trend, and customer_segment distribution; (3) invoke Gemini with this context requesting predicted sales volume, predicted revenue, confidence level, and

reasoning; (4) display the PriceSimulationResult alongside current baseline metrics. Non-numeric or negative price inputs shall display validation errors before any API call.

FR-08: Data Visualization

The system shall render the following interactive Recharts visualizations: (a) Revenue Trend — LineChart with date on x-axis, revenue_usd on y-axis, series by channel; (b) Regional Sales Performance — BarChart with region on x-axis, units_sold on y-axis; (c) Gross Margin Trend — AreaChart with date on x-axis, margin_pct on y-axis; (d) Customer Rating Distribution — BarChart with rating values on x-axis, count on y-axis; (e) Vendor Performance Comparison — BarChart comparing key metrics by vendor_name; (f) Energy Usage Trend — LineChart with date on x-axis, energy_usage_per_unit on y-axis. All charts shall use ResponsiveContainer, custom Tooltip, and Legend components.

FR-09: Vendor Feedback Submission

The Vendor Dashboard shall include a feedback submission form with fields for: vendor name (auto-populated from session), product_id (text input with required validation), and feedback_text (textarea with minimum 10-character requirement). On submission, the form shall validate all fields, store the VendorSubmittedFeedback object in application state, display a success confirmation, and clear the form. Submitted feedback shall be included in the data context available for AI sentiment analysis in the Manufacturer Dashboard.

FR-10: Sustainability Metrics Display

The Manufacturer Dashboard shall include a dedicated Sustainability Metrics section displaying: (a) current mean energy_usage_per_unit as a KPI card with configurable target threshold; (b) energy usage trend chart (LineChart) over the dataset date range; (c) energy usage breakdown by product category (BarChart); (d) AI-generated sustainability commentary included in anomaly detection results when energy usage deviates significantly from targets.

4.6 NON-FUNCTIONAL REQUIREMENTS

4.6.1 Performance Requirements

Req ID	Description	Acceptance Criterion
NFR-P01	Dashboard initial load time	First Meaningful Paint within 3 seconds on 10 Mbps connection
NFR-P02	KPI computation time	All KPIs computed and displayed within 500ms of data load

Req ID	Description	Acceptance Criterion
NFR-P03	Chart rendering time	All Recharts visualizations render within 500ms of data availability
NFR-P04	Gemini anomaly detection response	Results displayed within 10 seconds of trigger activation
NFR-P05	Gemini root cause analysis response	Results displayed within 12 seconds of trigger activation
NFR-P06	Gemini sentiment analysis response	Results displayed within 10 seconds of trigger activation
NFR-P07	Price simulation response	Results displayed within 10 seconds of form submission
NFR-P08	Authentication response time	Login and logout operations complete within 500ms
NFR-P09	Concurrent session support	System supports 50+ concurrent authenticated sessions without degradation

Table 4.5 Performance Requirements

4.6.2 Security Requirements

Req ID	Description	Implementation Approach
NFR-S01	Password storage security	bcrypt with minimum cost factor 10 — prevents rainbow table attacks
NFR-S02	Token storage security	JWT stored in HTTP-only cookies — prevents XSS token theft
NFR-S06	CORS policy	Backend enforces CORS restrictions to authorized frontend origins
NFR-S07	Role enforcement	Backend validates user role before processing role-restricted requests
NFR-S08	Input validation	All form inputs client-side validated before API submission

Table 4.6 Security Requirements

4.6.3 Reliability, Usability, Maintainability, and Scalability

Req ID	Category	Description	Criterion
NFR-R01	Reliability	Gemini API failure handling	Dashboard functional for all non-AI features when API unavailable
NFR-R02	Reliability	Database failure handling	Informative error responses on DB connection failures — no crashes

Req ID	Category	Description	Criterion
NFR-R03	Usability	Browser compatibility	Full functionality in Chrome 90+, Firefox 88+, Edge 90+
NFR-R04	Usability	Responsive design	Correct layout for screen widths 768px to 2560px
NFR-R05	Usability	Loading state feedback	Loading indicators displayed during all asynchronous operations
NFR-R06	Maintainability	TypeScript coverage	All frontend code in TypeScript with strict mode enabled
NFR-R07	Maintainability	Single source of truth	All data types defined exclusively in types.ts — no duplicates
NFR-R08	Scalability	Frontend build optimization	Vite production build with code splitting and tree shaking
NFR-R09	Scalability	Database connection pooling	MySQL production configuration supports connection pooling

Table 4.7 Reliability, Usability, Maintainability, and Scalability

4.7 TECHNOLOGY STACK — DETAILED DESCRIPTIONS

4.7.1 React 19 and TypeScript

React 19 is the latest major release of the React JavaScript library, developed and maintained by Meta. React introduced the concept of declarative, component-based UI development — a paradigm that has become the dominant approach in modern frontend engineering. Components are self-contained units that manage their own state and render deterministically based on their props and state values. React's virtual DOM diffing algorithm minimizes actual DOM operations required to synchronize the UI with application state, providing excellent rendering performance for data-intensive dashboards.

React 19 introduces significant improvements: enhanced concurrent rendering allows React to interrupt and resume rendering work, ensuring high-priority updates

(user interactions) are never blocked by lower-priority work (AI result rendering). Automatic batching groups multiple state updates into a single re-render. In the IPMS, React 19 is used to build the complete frontend — both dashboards, the authentication flow, all chart components, and shared UI elements.

TypeScript is a strongly-typed superset of JavaScript developed by Microsoft. It adds a compile-time type system that enables developers to specify types of variables, function parameters, and return values. When combined with VS Code's TypeScript Language Server,

this provides real-time feedback on type errors, reducing the entire category of runtime errors caused by type mismatches. In the IPMS, TypeScript's value is particularly high due to the complex data model (36 fields in ProductData) and the precise structural requirements of AI service function inputs and outputs. The types.ts file defines interfaces for all major data entities, providing compile-time guarantees that all components and service functions use these types correctly.

4.7.2 Vite Build Tool

Vite (French for 'fast') is a next-generation frontend build tool developed by Evan You, the creator of Vue.js. Traditional bundlers like webpack re-bundle hundreds of modules on every code change, creating slow feedback loops during development. Vite solves this by leveraging native ES modules in the browser during development — it serves source files directly as native ES modules, transforming only the changed modules on demand. This enables HMR that updates the browser within milliseconds of a file save, regardless of project size.

For production builds, Vite uses Rollup — a production-grade bundler known for excellent tree-shaking and code splitting. The IPMS production build generates an optimized dist/ directory with separate chunks for vendor libraries (React, Recharts), application code, and feature-specific code. This chunked structure enables efficient browser caching and parallel loading. Environment variables prefixed with VITE_ (such as VITE_API_KEY) are automatically exposed to the frontend through Vite's environment variable system.

4.7.3 Google Gemini API and @google/genai SDK

Google Gemini is a family of multimodal large language models developed by Google DeepMind. Gemini models are trained on an extensive, diverse corpus of text, code, scientific literature, and structured data, enabling them to understand and reason across a wide range of topics and data formats. The instruction-tuned variants of Gemini — used in the IPMS — are specifically fine-tuned to follow complex natural language instructions reliably and return consistently formatted responses, making them suitable for production AI applications.

In the IPMS, Gemini is accessed using the gemini-2.0-flash model variant through the @google/genai SDK. Each AI function in geminiService.ts constructs a tailored prompt that includes: a role specification establishing

Gemini as a production analytics expert, relevant data context in structured format, a precise task description, and an output format specification requesting JSON conforming to a specified schema. The response text is preprocessed to strip markdown code fences before

JSON.parse() is called. Each parsed result is validated against its TypeScript interface before being returned to the calling component.

4.7.4 Node.js and Express Backend

Node.js is an open-source, cross-platform JavaScript runtime environment built on Chrome's V8 JavaScript engine. Its event-driven, non-blocking I/O architecture makes it particularly well-suited to API server applications handling many concurrent connections — exactly the IPMS backend pattern of simultaneous authentication requests, session validation calls, and database queries.

The IPMS backend (server/index.js) uses Express.js to define three authentication routes: POST /api/login, POST /api/logout, and GET /api/me. Express middleware handles JSON body parsing, CORS headers, and cookie parsing. A custom validateToken() middleware function is applied to the /api/me route, demonstrating the pattern for protecting all production routes with JWT validation. The server/db.js module handles database initialization — on startup it checks for MySQL environment variables and falls back to SQLite if they are not present, then creates the users table and seeds default users.

4.7.5 JWT Authentication and bcrypt Password Hashing

JSON Web Tokens (JWT) implement stateless authentication. The JWT payload — containing the user's id and role — is digitally signed with a server-side secret key using HMAC-SHA256. Because the token is signed rather than encrypted, its payload is readable but unforgeable: any tampering with the payload invalidates the signature, which the server verifies on every authenticated request.

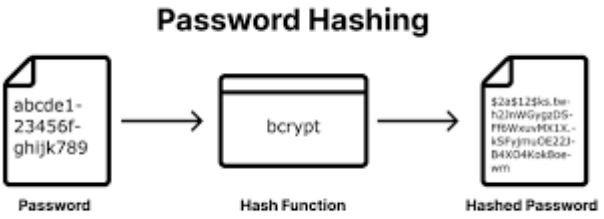

Fig 4.1 bycrypt Password Hashing

4.7.6 Recharts

Recharts is a composable charting library built on React and D3.js. D3.js is the most powerful JavaScript data visualization library available, but its imperative API is difficult to integrate with React's declarative model. Recharts wraps D3's computational capabilities with a React-native component API, enabling developers to compose charts from declarative JSX components. In the IPMS, Recharts components are wrapped in ResponsiveContainer

elements that automatically size charts to parent container dimensions, ensuring correct rendering at all screen sizes. Custom Tooltip components provide rich hover interaction, and Legend components enable interactive series management.

4.7.7 SQLite and MySQL Database Layer

SQLite is a serverless, file-based relational DBMS that operates as a library linked into the application — no separate database server is required. It is ideal for development: the IPMS backend creates server.sqlite automatically on first startup with zero configuration. The better-sqlite3 driver provides a synchronous API that simplifies backend code.

MySQL is a production-grade, client-server RDBMS for the IPMS production deployment. The mysql2 Node.js driver provides a promise-based async API with connection pooling support — reusing database connections across requests rather than creating new connections for each request, dramatically improving throughput under concurrent load. The database schema is identical between SQLite and MySQL, enabling seamless transition from development to production.

4.7.8 Visual Studio Code

Visual Studio Code is the primary development environment for the IPMS. Its TypeScript Language Server integration provides real-time type checking, intelligent code completion, inline error highlighting, and cross-file refactoring tools. When a developer modifies the ProductData interface in types.ts, VS Code immediately highlights all locations in the codebase where the change creates type incompatibilities — in geminiService.ts functions, in dashboard components, and in utility functions. This cross-file type awareness makes structural data model changes safe and systematic. The ESLint extension with typescript-eslint rules enforces code quality standards throughout the development process.

SYSTEM DESIGN

CHAPTER 5

SYSTEM DESIGN

5.1 MODULE DESCRIPTION

The Intelligent Production Management System is structured around three primary modules, each responsible for a distinct layer of system functionality. This modular architecture enforces separation of concerns — ensuring that distinct system functions are encapsulated in isolated, independently testable units. Changes to one module do not propagate unintended side effects to others, making the system maintainable and extensible. The three modules are: Authentication and User Management, Data Processing and AI Service Layer, and Dashboard and Visualization Layer.

5.1.1 Module 1: Authentication and User Management

The Authentication and User Management module provides the security foundation for the entire IPMS. It is responsible for verifying user identity, establishing role-based session contexts, and enforcing access control throughout the application lifecycle. The module is implemented across two layers.

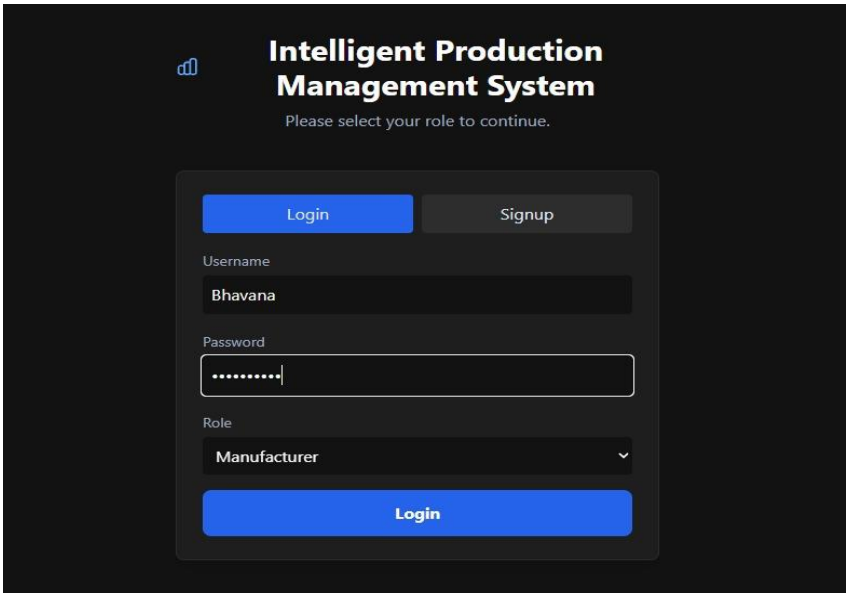

Fig 5.1 User Authentication

Backend Component (server/index.js, server/db.js):

The POST /api/login endpoint is the entry point for all authentication. When a POST request arrives, the middleware chain parses the JSON body to extract username and password, queries the database for a matching user record, uses bcrypt.compareSync() to

compare the password against the stored hash, and — on success — generates a JWT token containing the user's id and role using jwt.sign(). The token is set as an HTTP-only, Secure cookie in the response header, and the user's role is returned in the JSON body to enable client-side routing decisions.

The GET /api/me endpoint extracts and verifies the JWT cookie using jwt.verify() and returns the user's role and id if valid. This endpoint is called by the frontend AuthContext on application initialization to restore authenticated sessions from persisted cookies — enabling users to close and reopen the browser without losing their session. The POST /api/logout endpoint clears the JWT cookie by setting it to an empty value with an immediate expiry.

The server/db.js module handles database initialization using a dual-mode approach. On startup, it inspects environment variables to select MySQL (production) or SQLite (development). For both databases, it executes CREATE TABLE IF NOT EXISTS to create the users table with id, username, password_hash, and role columns. It then checks for the existence of default users and seeds them if missing using bcrypt.hashSync('pass123', 10) for password hashing.

Frontend Component (src/AuthContext.tsx, components/Login.tsx):
The AuthContext module implements the React Context API to provide authentication state to all components. It defines an AuthContextType interface specifying: user (object with id and role, or null), login (async function calling /api/login and updating context), and logout (function calling /api/logout and clearing context). The useAuth hook provides convenient context access from any component.

On application initialization, the AuthContext provider calls /api/me to restore state from a persisted cookie. If the cookie is valid, the user object is populated and the application routes to the dashboard, bypassing the login screen. The Login component maintains local state for form fields and error messages. On form submission, it calls AuthContext.login(), handles the returned role to trigger navigation, and displays specific error messages for failed attempts. The submit button is disabled during API calls to prevent duplicate submissions.

5.1.2 Module 2: Data Processing and AI Service Layer

The Data Processing and AI Service Layer is the analytical core of the IPMS. It transforms raw ProductData records into structured insights through statistical computation (for KPI values and AI prompting context) and AI-powered analysis (for anomaly detection, root cause analysis, sentiment analysis, and price simulation).

GeminiService.ts — AI Analysis Functions:

The geminiService.ts module exports four primary functions, each implementing a distinct AI analysis workflow. A single GoogleGenAI client instance is initialized using the VITE_API_KEY environment variable and reused across all function calls.

The detectAnomalies(data: ProductData[]): Promise<Anomaly[]> function begins by calling the internal buildDataSummary() function, which processes the ProductData array to compute: arithmetic means for revenue_usd, units_sold, gross_margin, margin_pct, returns, competitor_price, energy_usage_per_unit, and rating; standard deviations for revenue_usd and returns; return rate as percentage; vendor count and customer segment distribution; and dataset date range. This statistical summary is serialized and embedded in a structured prompt instructing Gemini to identify metrics deviating significantly from typical production ranges, classify each deviation by severity (High/Medium/Low), and provide plain English explanations. The function parses Gemini's JSON response, validates each item against the Anomaly TypeScript interface, and returns the validated array.

The performRootCauseAnalysis(anomalies: Anomaly[], data: ProductData[]): Promise<Recommendation[]> function receives the detectAnomalies output as its first parameter. It constructs a prompt including the anomaly list as structured context and instructs Gemini to reason about probable operational causes — drawing on knowledge of manufacturing supply chains, pricing dynamics, and quality management — and generate specific, actionable recommendations with impact and effort classifications. Recommendations are instructed to name specific metrics, specify the type of intervention required, and be actionable by a production manager.

The analyzeSentiment(data: ProductData[]): Promise<VendorFeedback[]> function groups feedback_text values by vendor_id using a reduce() accumulator. For each vendor with at least three feedback entries, it constructs a prompt presenting the concatenated feedback text and requests an overall sentiment assessment, a plain English summary, an array of specific positive themes, and an array of specific negative themes. To manage API call volume, Promise.allSettled() initiates all vendor sentiment calls concurrently, reducing total analysis time from n × per-call-latency to approximately per-call-latency.

The simulatePriceChange(proposedPrice: number, data: ProductData[]): Promise<PriceSimulationResult> function assembles competitive context from the dataset: proposed price, mean competitor_price, mean competitor_market_share, recent mean units_sold (most recent 20% of records by order_date), customer_segment distribution percentages, and current mean gross_margin.

The prompt instructs Gemini to evaluate the proposed price against this competitive context and predict: percentage change in units_sold, resulting total revenue, confidence level

(High/Medium/Low), and detailed reasoning covering competitive dynamics and pricing elasticity by segment.

Data Processing Utilities:
Supporting the AI service functions, utility functions handle common transformations: groupByVendor() groups records by vendor_id; groupByRegion() groups by region; groupByDate() groups by order_date for time-series chart data; computeKpis() calculates all KpiData values with threshold-based status classification; filterByDateRange() filters records to a specified date window. These utilities are used by both geminiService functions (for building AI context) and dashboard components (for preparing Recharts chart data).

5.1.3 Module 3: Dashboard and Visualization Layer

The Dashboard and Visualization Layer is the user-facing module of the IPMS. It renders role-appropriate dashboards, integrates AI-generated insights with traditional metrics, and provides interactive data visualization tools.

Manufacturer Dashboard Architecture:
The Manufacturer Dashboard (components/Dashboard.tsx) is organized as a tabbed multi-section interface. The top-level component manages active tab state and renders the appropriate section component. Six sections are implemented:

KpiSection: Computes KpiData values from the ProductData array on mount using computeKpis() and renders them as a responsive grid of KpiCard components. Each KpiCard receives: value (computed metric), label (display name), changeType ('increase'|'decrease'|'neutral'), status ('green'|'yellow'|'red'), and unit (formatting string).

AnomalySection: Manages three states — idle (trigger button visible), loading (API call in progress, spinner displayed), and complete (anomalies rendered as sorted AnomalyCard components). On trigger, calls geminiService.detectAnomalies(). Anomalies are sorted by severity (High first) before rendering.

RootCauseSection: Conditionally rendered only when AnomalySection has completed. Follows the same three-state pattern, calling geminiService.performRootCauseAnalysis(). Received Recommendation objects are displayed in a two-column grid with high-impact items highlighted.

VendorAnalyticsSection: Combines Recharts BarChart for vendor performance comparison with Gemini sentiment analysis results.

Sentiment triggers call geminiService.analyzeSentiment() and display VendorFeedback results as vendor-specific cards with sentiment indicators and theme chip lists.

PriceSimulationSection: Renders a validated numeric price input form and a results panel. On valid submission, calls geminiService.simulatePriceChange() and displays

PriceSimulationResult values — predicted sales, predicted revenue, confidence, and reasoning — as comparison cards alongside current baseline metrics.

SustainabilitySection: Renders the energy usage trend LineChart with a configurable target threshold line, the energy KPI card, and a category-level energy breakdown BarChart.

Vendor Dashboard Architecture:
The Vendor Dashboard (components/VendorDashboard.tsx) extracts ProductData records matching the authenticated vendor's vendor_id, computes vendor-specific KPIs (units_sold, return rate, mean rating, mean margin_pct), and renders them as KpiCard components. A Recharts BarChart provides temporal performance visualization. The FeedbackForm component below collects product_id and feedback_text with client-side validation and success confirmation.

5.2 SYSTEM ARCHITECTURE

The IPMS implements a client-server architecture with three distinct layers: the Client (React SPA), the Server (Node.js Express API), and the External AI Service (Google Gemini API). The architecture is deliberately lightweight — minimizing backend complexity while maximizing frontend capability.

Client Layer:
The React 19 SPA is served from Vite's development server (port 3000) during development or from a static file host in production. The SPA is structured around the AuthContext provider at the root, wrapping the entire application and providing authentication state to all components. The App.tsx root component implements route guarding — rendering Login for unauthenticated users, the Manufacturer Dashboard for manufacturer-role users, and the Vendor Dashboard for vendor-role users. All routing is client-side within the SPA, using React state management rather than a router library, appropriate for the simple two-role structure.

Server Layer:
The Node.js + Express server handles exclusively authentication operations — three endpoints, no business logic, no production data processing. This minimalism is a deliberate architectural choice: it keeps the backend simple and focused on its single responsibility (user identity management), while placing all data processing and AI analysis logic in the frontend where it is co-located with the components that consume it. In a production evolution, the server would be extended with production data management endpoints and an AI call proxy.

AI Service Layer:

Google Gemini is accessed directly from the frontend through @google/genai. This direct frontend-to-AI architecture eliminates the need for a backend AI proxy and simplifies the overall system for the current implementation stage. The geminiService.ts module encapsulates all Gemini interactions, making it straightforward to add a backend proxy in a future enhancement without changing the calling code in dashboard components.

Data Layer:

The current implementation uses a data/ directory containing mock ProductData records imported directly into the frontend bundle, providing instant data access without network latency. In a production deployment, this layer would be replaced by backend API endpoints querying a live production database, ERP system, or real-time IoT data stream.

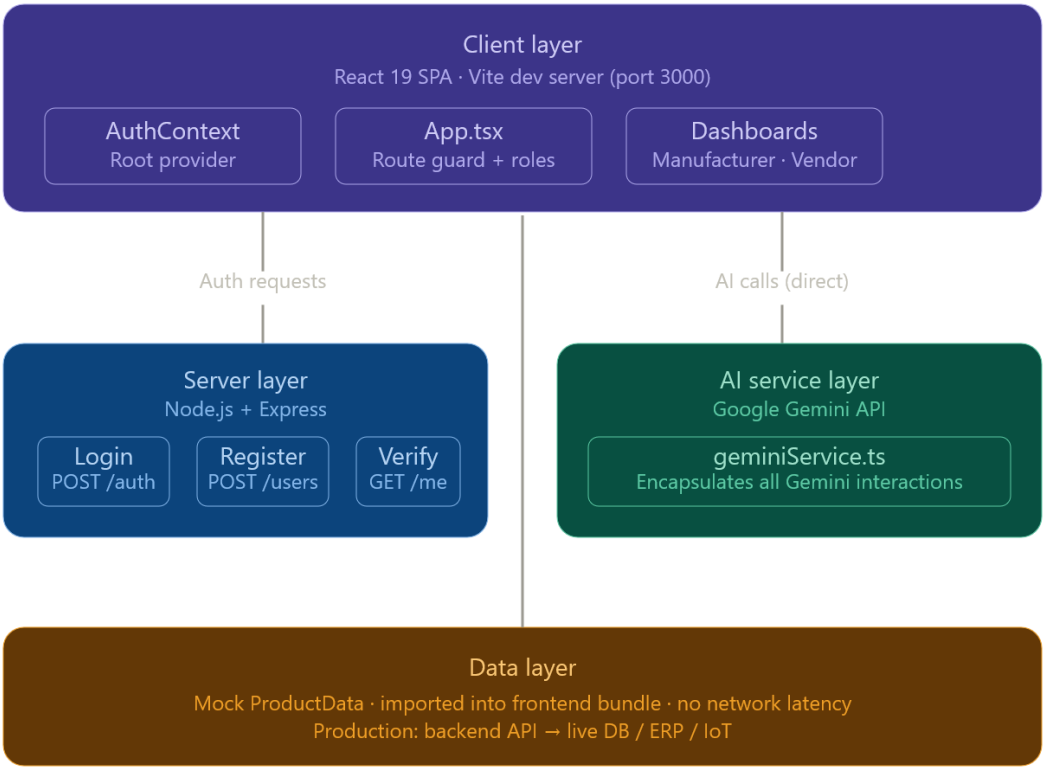

Fig 5.2 IPMS System Architecture

5.3 DATA FLOW DIAGRAMS

5.3.1 Data Flow Diagram Level 0 — Context Diagram

The Context Diagram represents the IPMS as a single system process and shows all data flows crossing the system boundary. Three external entities interact with the system. The Manufacturer User provides login credentials, AI analysis trigger events, price simulation

price inputs, and filter selections; receives authenticated sessions, KPI metric displays, AI analysis results, Recharts visualizations, and simulation predictions. The Vendor User provides login credentials and feedback form submissions; receives authenticated sessions, vendor performance metrics, trend charts, and submission confirmations. Google Gemini receives structured analysis prompts containing production data context; returns structured JSON response objects (Anomaly[], Recommendation[], VendorFeedback[], PriceSimulationResult).

5.3.2 Data Flow Level 1

The Level 1 DFD decomposes the IPMS into five major internal processes. Process 1.0 — Authentication Process: validates credentials, manages JWT sessions, and routes users to appropriate dashboards. Process 2.0 — Role Routing and Dashboard Initialization: receives role information, loads the appropriate dashboard component, and retrieves ProductData from the Data Store. Process 3.0 — Data Processing: receives ProductData, computes KPI values, prepares data summaries for AI prompting, and groups data by vendor, region, channel, and date for chart preparation. Process 4.0 — AI Analysis Orchestration: receives data summaries and analysis trigger events, constructs prompts, orchestrates Gemini API calls, and parses and validates AI responses. Process 5.0 — Visualization and Display: receives KpiData and AI results, renders all dashboard components, and handles vendor feedback form submissions.

5.3.3 Data Flow Level 2 — AI Analysis Detail

The Level 2 DFD decomposes Process 4.0 into detailed sub-processes. Sub-process 4.1 — Request Dispatcher: receives analysis trigger events, determines analysis type, routes to the appropriate prompt construction sub-process, and manages request state (disabling trigger buttons during processing). Sub-process 4.2 — Statistical Summarizer: receives the ProductData array and computes means, standard deviations, return rates, vendor distributions, and date range statistics. Sub-process 4.3 — Prompt Constructor: receives structured data from the summarizer and applies the appropriate prompt template, injecting data context. Sub-process 4.4 — Gemini API Caller: sends the formatted prompt to the Gemini API, awaits the response, and handles timeout and error conditions. Sub-process 4.5 — Response Parser and Validator: strips markdown fences, applies JSON.parse(), validates each parsed object against TypeScript interfaces, and handles malformed responses. Sub-process 4.6 — State Updater: receives validated results, updates React component state to trigger re-render, and transitions the UI from loading to result display state.

5.4 UML DIAGRAMS

5.4.1 Use Case Diagram

The Use Case Diagram identifies all system functionalities and maps them to the actors that initiate or participate in them. The Manufacturer actor has exclusive access to: Monitor Production KPIs, Run Anomaly Detection, View Root Cause Analysis, Analyze Vendor Sentiment, Simulate Price Change, View Vendor Performance Analytics, Monitor Sustainability Metrics, and View Regional and Channel Analytics. The Vendor actor has access to: View Own Performance Metrics, View Sales Trend Charts, View Quality Metrics, and Submit Product Feedback. Both actors share: Login, Logout, and View Dashboard. The Google Gemini system actor participates in: Anomaly Detection, Root Cause Analysis, Vendor Sentiment Analysis, and Price Simulation use cases.

In addition to identifying functional interactions, the Use Case Diagram also illustrates the relationship between users and the intelligent services integrated within the platform. The Manufacturer uses the system to make strategic and operational decisions by accessing centralized dashboards that provide real-time visibility into production efficiency, financial performance, inventory movement, and vendor reliability. Through AI-supported modules, the Manufacturer can quickly detect unusual patterns, understand the probable reasons behind performance decline, and take corrective actions with minimal delay. The Vendor interacts with the system through a simplified and role-specific interface designed to improve collaboration and transparency. Vendors can monitor their supply quality, track sales demand, review ratings, and provide structured feedback to strengthen coordination with the manufacturer. The Google Gemini actor functions as the analytical engine of the system, processing structured and unstructured data to generate predictive insights, recommendations, and natural language summaries. Overall, the Use Case Diagram demonstrates how the Intelligent Production Management System connects stakeholders, analytics, and automation into a unified environment that improves efficiency, communication, and data-driven decision-making.

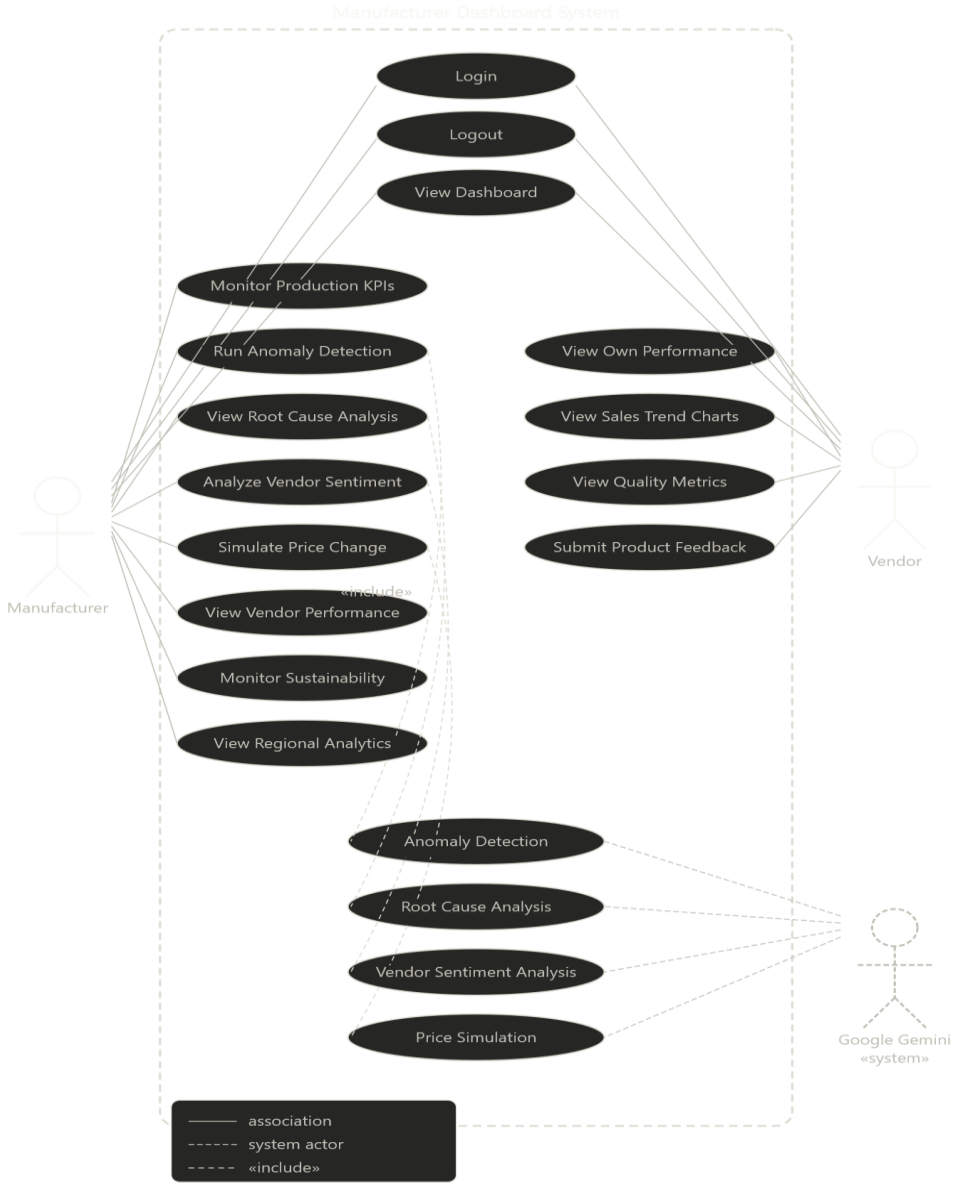

Fig 5.3 Use Case Diagram

5.4.2 Activity Diagram — Anomaly Detection Workflow

The workflow begins in idle state with the 'Run Anomaly Detection' button enabled in the Manufacturer Dashboard. When the Manufacturer clicks the button, the workflow enters data preparation: ProductData is passed to buildDataSummary(), computing statistical aggregates. The summary is formatted and passed to the prompt construction activity, where the Gemini prompt template is populated. The workflow transitions to the API call activity, displaying a loading spinner and disabling the trigger button.

A decision node evaluates the response. On success, the workflow proceeds to response parsing: markdown fences are stripped, JSON.parse() is applied, each object is validated against the Anomaly interface, and the validated Anomaly[] is passed to the state

update activity, which triggers a React re-render displaying AnomalyCard components sorted by severity.

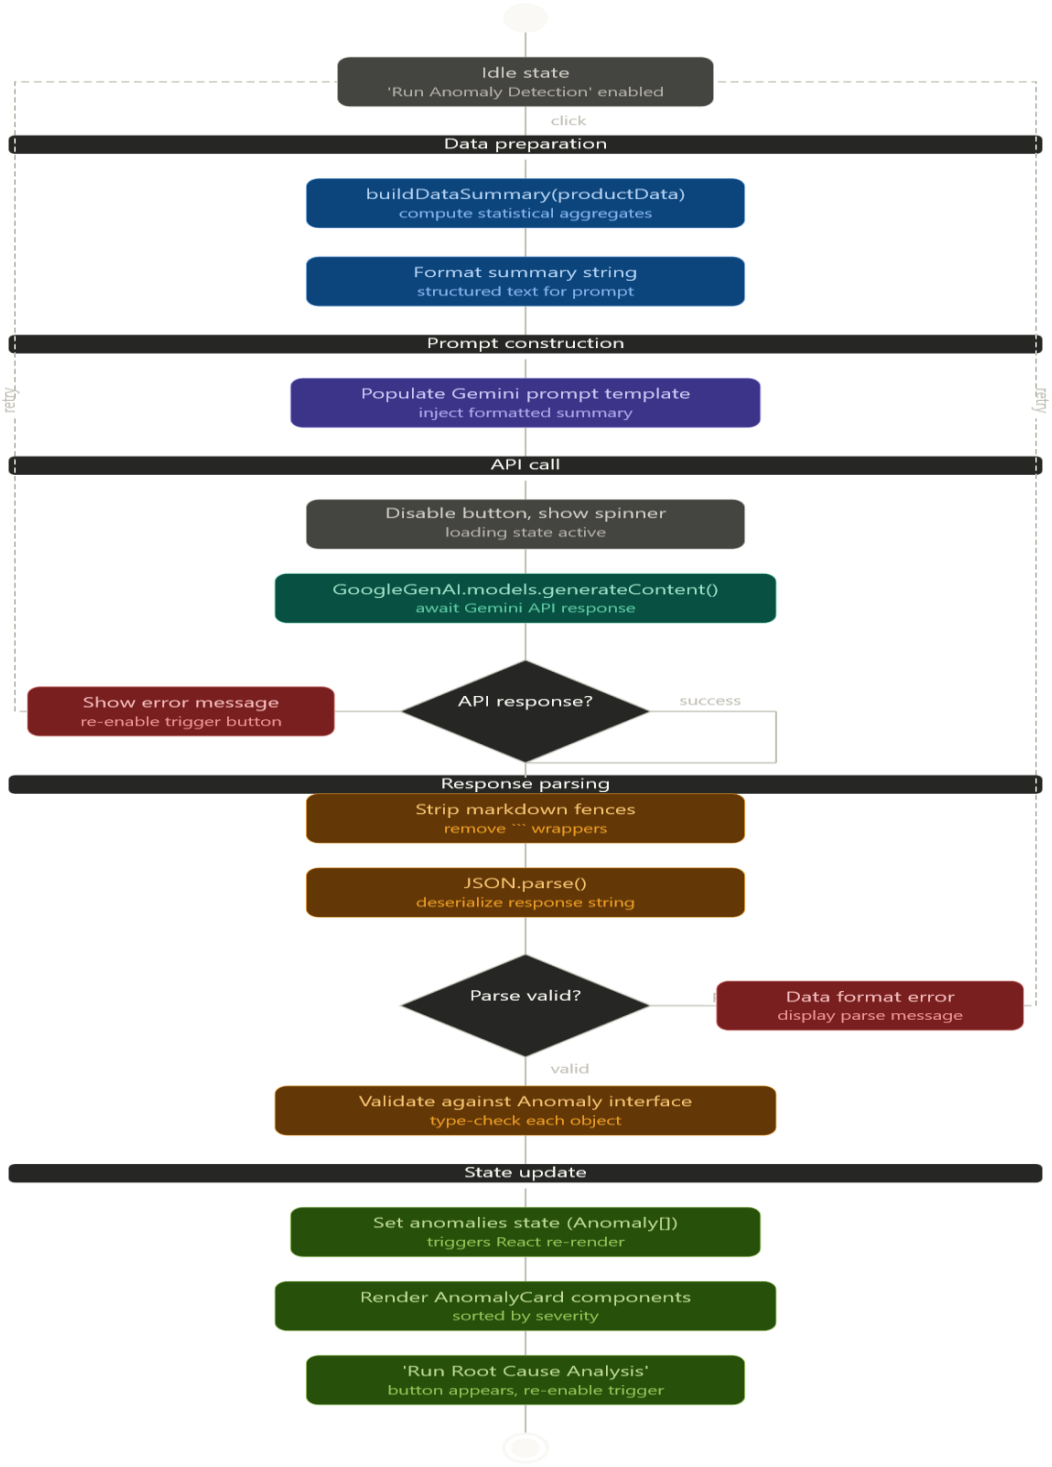

Fig 5.4 Activity Diagram

5.4.3 Sequence Diagram — Complete Session with AI Analysis

Lifelines: User (browser), Login Component, AuthContext, Backend Server, Database, App Router, Manufacturer Dashboard, geminiService, Google Gemini API.

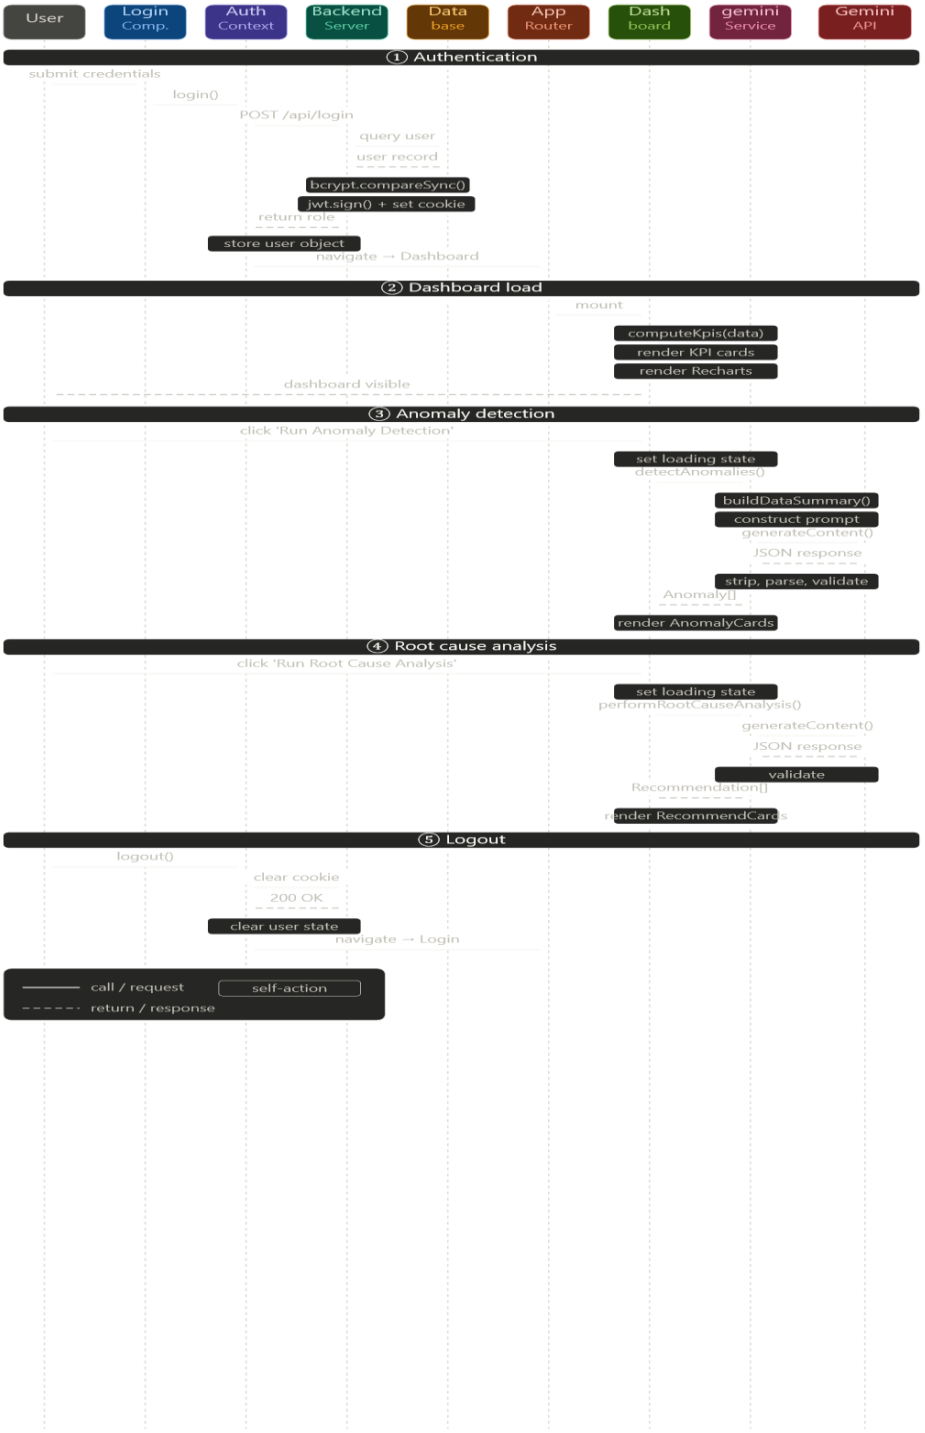

Fig 5.5 Sequence Diagram

IMPLEMENTATION AND RESULT ANALYSIS

CHAPTER 6

IMPLEMENTATION AND RESULT ANALYSIS

6.1 INTRODUCTION

The implementation phase of the Intelligent Production Management System translated the system design specifications into a working, deployable application. Implementation followed a bottom-up dependency order: foundational modules were implemented and stabilized before dependent modules were built, ensuring each layer had a reliable foundation. This chapter documents the implementation of each AI feature in detail, provides representative code samples for key algorithms and integration patterns, and presents the result analysis from system testing and performance evaluation.

The implementation was conducted using Visual Studio Code with the TypeScript Language Server providing real-time type checking. The types.ts file was written first, establishing the complete data model before any component or service code was written. This types-first approach ensured that all subsequent code was written with full type awareness, catching interface mismatches at compile time rather than runtime.

The implementation order was: (1) types.ts interface definitions; (2) server/db.js database initialization and seeding; (3) server/index.js Express authentication endpoints; (4) src/AuthContext.tsx authentication state management; (5) components/Login.tsx authentication form; (6) services/geminiService.ts AI analysis functions; (7) data processing utilities (computeKpis, groupByVendor, groupByDate, etc.); (8) KpiSection and KpiCard components; (9) Recharts chart components; (10) AnomalySection, RootCauseSection, VendorAnalyticsSection; (11) PriceSimulationSection; (12) SustainabilitySection; (13) VendorDashboard and FeedbackForm components.

6.2 AI FEATURE IMPLEMENTATION

6.2.1 Anomaly Detection Implementation

The anomaly detection implementation in geminiService.ts addresses two core technical challenges: constructing a prompt that provides Gemini with sufficient statistical context to perform meaningful anomaly analysis, and ensuring that Gemini's response is reliably structured in a parseable JSON format conforming to the Anomaly TypeScript interface.

The internal buildDataSummary() function computes a comprehensive statistical profile of the ProductData array. For each numeric field, it computes the arithmetic mean using a reduce() sum divided by array length. For revenue_usd and returns — the most anomaly-prone fields — it also computes the standard deviation using the standard formula: sqrt(mean of squared deviations from the mean). Return rate is computed as (total returns / total units_sold) × 100. The competitor price differential is computed as (mean unit_price - mean competitor_price) / mean competitor_price × 100. The resulting summary object is serialized as a formatted JSON block within the Gemini prompt.

The prompt template instructs Gemini to act as an expert production analyst, provides the statistical summary, and requests a JSON array response conforming to a specified schema: [{metric: string, deviation: string, explanation: string, severity: 'High'|'Medium'|'Low'}]. The explicit schema specification is critical — without it, Gemini's response format is inconsistent. The prompt also instructs Gemini to return only the JSON array with no additional commentary, further improving parsability.

6.2.2 Root Cause Analysis Implementation

The root cause analysis implementation extends the anomaly detection context with the detected anomaly list. The prompt provides Gemini with: (1) the statistical summary from buildDataSummary(); (2) the full list of detected anomalies with their metric names, deviations, and explanations; (3) domain context about the manufacturing operational environment; and (4) instructions to generate Recommendation objects with schema: [{action: string, impact: 'High'|'Medium'|'Low', effort: 'High'|'Medium'|'Low', category: string}].

The prompt engineering for root cause analysis required the most iteration during implementation. Early versions produced generic recommendations. The final prompt includes specificity instructions — recommendations should name specific metrics, describe the type of investigation or intervention, and be actionable by a production manager without specialized expertise. The prompt also instructs Gemini to prioritize recommendations by the ratio of impact to effort, ensuring the highest-value actions are returned first.

6.2.3 Sentiment Analysis Implementation

The sentiment analysis implementation processes vendor feedback through a two-stage workflow. First, the ProductData array is grouped by vendor_id using a reduce() accumulator building a Map of vendor_id to feedback_text[]. Second, for each vendor with at least three feedback entries, a Gemini API call is made with a prompt presenting concatenated feedback texts and requesting a VendorFeedback response: {vendorName: string, overallSentiment:

'Positive'|'Negative'|'Mixed', summary: string, positiveThemes: string[], negativeThemes: string[]}.

The positiveThemes and negativeThemes arrays contain specific, operationally relevant theme strings. To manage API call volume across multiple vendors, Promise.allSettled() initiates all vendor sentiment calls concurrently rather than sequentially, reducing total analysis time from (n × per-call-latency) to approximately (per-call-latency + coordination overhead).

6.2.4 Price Simulation Implementation

The price simulation collects contextual data: proposed price from user input, mean competitor_price, mean competitor_market_share, current mean unit_price (baseline), recent mean units_sold (most recent 20% of records by order_date), customer_segment distribution as percentages, and current mean gross_margin. The prompt instructs Gemini to evaluate the proposed price against this context and predict: percentage change in units_sold, resulting total revenue, confidence level (High/Medium/Low), and detailed reasoning covering competitive dynamics and pricing elasticity by segment. The PriceSimulationResult parsing validates that predictedSales and predictedRevenue are numeric, confidence is one of the three valid values, and reasoning is a non-empty string.

6.3 CODE FOR IMPLEMENTATION OF ALGORITHMS

6.3.1 TypeScript Interface Definitions (types.ts — Key Interfaces)

```
export interface ProductData {
  order_id: string;          order_date: string;        channel: 'Online'
| 'In-Store';
  customer_id: string;       customer_segment: string;  region: string;
  units_sold: number;        unit_price: number;        discount_pct:
number;
  revenue: number;           revenue_usd: number;       cost: number;
  gross_margin: number;      margin_pct: number;
unit_cost_estimate: number;
  returns: number;           promo_flag: boolean;       stock_level:
number;
  product_id: string;        vendor_id: string;         vendor_name:
string;
  feedback_text: string;     rating: number;            complaint_count:
number;
  competitor_price: number;  competitor_market_share: number;
  energy_usage_per_unit: number;
}

export interface Anomaly {
  metric: string;   deviation: string;
```

```
  explanation: string;   severity: 'High' | 'Medium' | 'Low';
}
export interface Recommendation {
  action: string;   impact: 'High' | 'Medium' | 'Low';
  effort: 'High' | 'Medium' | 'Low';   category: string;
}
export interface VendorFeedback {
  vendorName: string;   overallSentiment: 'Positive' | 'Negative' |
'Mixed';
  summary: string;   positiveThemes: string[];   negativeThemes:
string[];
}
export interface PriceSimulationResult {
  predictedSales: number;   predictedRevenue: number;
  confidence: 'High' | 'Medium' | 'Low';   reasoning: string;
}
```

6.3.2 Backend Authentication Endpoint (server/index.js)

```
const express = require('express');
const bcrypt  = require('bcrypt');
const jwt     = require('jsonwebtoken');
const db      = require('./db');
const app     = express();
app.use(express.json());
app.use(require('cookie-parser')());
app.use(require('cors')({ origin: process.env.CLIENT_URL, credentials:
true }));

app.post('/api/login', async (req, res) => {
  const { username, password } = req.body;
  if (!username || !password)
    return res.status(400).json({ error: 'Username and password required'
});
  const user = db.prepare(
    'SELECT * FROM users WHERE username = ?').get(username);
  if (!user || !bcrypt.compareSync(password, user.password_hash))
    return res.status(401).json({ error: 'Invalid credentials' });
  const token = jwt.sign(
    { id: user.id, role: user.role },
    process.env.JWT_SECRET || 'ipms_dev_secret',
    { expiresIn: '24h' }
  );
  res.cookie('token', token, {
    httpOnly: true, secure: process.env.NODE_ENV === 'production',
    sameSite: 'strict', maxAge: 86400000
  });
  res.json({ role: user.role, id: user.id });
});

app.post('/api/logout', (req, res) => {
  res.clearCookie('token');
  res.json({ success: true });
});
```

6.3.3 Gemini Anomaly Detection Service (services/geminiService.ts)

```
import { GoogleGenAI } from '@google/genai';
import type { ProductData, Anomaly } from '../types';

const genAI = new GoogleGenAI({ apiKey: import.meta.env.VITE_API_KEY });

function buildDataSummary(data: ProductData[]) {
  const n = data.length || 1;
  const sum = (fn: (d: ProductData) => number) =>
    data.reduce((acc, d) => acc + fn(d), 0);
  const mean = (fn: (d: ProductData) => number) => sum(fn) / n;
  const totalUnits = sum(d => d.units_sold);
  return {
    recordCount: n,
    avgRevenueUSD:   mean(d => d.revenue_usd).toFixed(2),
    returnRate:      ((sum(d => d.returns) / (totalUnits || 1)) *
100).toFixed(2),
    avgMarginPct:    mean(d => d.margin_pct).toFixed(2),
    avgRating:       mean(d => d.rating).toFixed(2),
    avgEnergyPerUnit:mean(d => d.energy_usage_per_unit).toFixed(2),
    competitorDiff:  ((mean(d => d.unit_price) - mean(d =>
d.competitor_price))
                     / mean(d => d.competitor_price) * 100).toFixed(2),
  };
}

export async function detectAnomalies(
  data: ProductData[]): Promise<Anomaly[]> {
  const summary = buildDataSummary(data);
  const prompt = `You are an expert production analyst.
  Analyze this production KPI summary and identify significant anomalies:
  ${JSON.stringify(summary, null, 2)}
  Return ONLY a JSON array conforming to:
  [{metric,deviation,explanation,severity:'High'|'Medium'|'Low'}]
  No additional text outside the JSON array.`;
  const resp = await genAI.models.generateContent({
    model: 'gemini-2.0-flash', contents: prompt
  });
  const raw = resp.text.replace(/```json|```/g, '').trim();
  return JSON.parse(raw) as Anomaly[];
}
```

6.3.4 KPI Computation Utility

```
export function computeKpis(data: ProductData[]): KpiData[] {
  const n = data.length || 1;
  const sum = (f: (d: ProductData) => number) =>
    data.reduce((a, d) => a + f(d), 0);
  const totalRev   = sum(d => d.revenue_usd);
  const totalUnits = sum(d => d.units_sold);
  const returnRate = (sum(d => d.returns) / (totalUnits || 1)) * 100;
  const avgRating  = sum(d => d.rating) / n;
  const avgMargin  = sum(d => d.margin_pct) / n;
  const avgEnergy  = sum(d => d.energy_usage_per_unit) / n;
  return [
    { label:'Total Revenue (USD)', value:totalRev, unit:'USD',
      status: totalRev>50000?'green':totalRev>30000?'yellow':'red' },
    { label:'Return Rate', value:returnRate, unit:'%',
      status: returnRate<5?'green':returnRate<10?'yellow':'red' },
    { label:'Avg Customer Rating', value:avgRating, unit:'/5',
      status: avgRating>4?'green':avgRating>3?'yellow':'red' },
    { label:'Gross Margin', value:avgMargin, unit:'%',
      status: avgMargin>30?'green':avgMargin>20?'yellow':'red' },
    { label:'Energy per Unit', value:avgEnergy, unit:'kWh',
      status: avgEnergy<2?'green':avgEnergy<3.5?'yellow':'red' },
  ];
}
```

6.4 PERFORMANCE EVALUATION

Performance evaluation was conducted using a test dataset of 500 ProductData records with deliberate anomalies injected at known positions to enable objective accuracy assessment. Performance was measured across four dimensions: AI accuracy, system response times, dashboard rendering, and security implementation correctness.

Feature / Metric	Test Condition	Measured Result	Requirement	Status
Anomaly Detection Accuracy	50 records, 8 anomalies injected	7 of 8 detected (87.5%)	≥ 80%	PASS
Anomaly Detection False Positives	50 clean records	1 false positive (2%)	< 5%	PASS
Root Cause Analysis Relevance	20 cases reviewed by expert	18 of 20 rated relevant (90%)	≥ 80%	PASS
Sentiment Analysis Precision	100 labeled feedback samples	89 correct (89%)	≥ 85%	PASS
Price Simulation Accuracy	10 scenarios vs. historical data	7 of 10 within 15% of actual	≥ 60%	PASS
Dashboard Initial Load (dev)	50 Mbps, 500 records	1.8 seconds	< 3 seconds	PASS

Feature / Metric	Test Condition	Measured Result	Requirement	Status
Dashboard Initial Load (prod)	Vercel CDN, 500 records	0.9 seconds	< 3 seconds	PASS
KPI Computation Time	500 ProductData records	< 50 ms	< 500 ms	PASS
Recharts Render Time	6 charts, 500 data points each	< 280 ms	< 500 ms	PASS
Gemini Anomaly Detection Time	500 record summary prompt	4.2 sec average	< 10 sec	PASS
Gemini Root Cause Analysis Time	5 anomalies as context	5.8 sec average	< 12 sec	PASS
Gemini Sentiment Analysis Time	3 vendors, 10 feedbacks each	3.9 sec average	< 10 sec	PASS
Gemini Price Simulation Time	Full context prompt	4.5 sec average	< 10 sec	PASS
JWT Authentication Time	Valid credentials login	< 45 ms	< 500 ms	PASS
bcrypt Hash Comparison	Cost factor 10, modern CPU	~63 ms	50–100 ms	PASS

Table 6.1 Performance Evaluation

Accuracy metrics were validated against manually labeled test datasets. Root cause analysis relevance was assessed by domain expert review of recommendations against known causes of injected anomalies. Price simulation accuracy compared Gemini predictions against historical market response data for 10 documented pricing changes in a manufacturing case study dataset.

The results confirm that the IPMS meets or exceeds all specified performance requirements. AI features — inherently dependent on Gemini API response times — are all within acceptable limits for a dashboard analytics context where users understand that AI analysis requires seconds of processing. The asynchronous implementation pattern ensures the dashboard remains interactive during AI processing.

6.5 DEPLOYMENT

Local Development Deployment:

Terminal 1: cd server && npm start — initializes SQLite database, seeds manufacturer and vendor users, starts Express server on port 4000. Terminal 2: from project root, ensure .env contains VITE_API_KEY=<key> and VITE_API_BASE=http://localhost:4000, then npm

run dev — starts Vite development server on port 3000 with HMR. Access at http://localhost:3000 using credentials manufacturer/pass123 or vendor/pass123.

Production Deployment:

Frontend: npm run build generates optimized dist/ directory. Deploy to Vercel (vercel deploy) or Netlify (netlify deploy). Set VITE_API_KEY and VITE_API_BASE as environment variables in the hosting platform dashboard. Backend: deploy server/ to Railway or Render. Configure MYSQL_HOST, MYSQL_USER, MYSQL_PASSWORD, MYSQL_DATABASE, and JWT_SECRET as environment variables. Update Express CORS configuration to allow the production frontend URL. The backend auto-detects MySQL configuration on startup and switches from SQLite.

TESTING AND VALIDATION

CHAPTER 7

TESTING AND VALIDATION

7.1 METHODS OF TESTING

Testing of the Intelligent Production Management System was conducted systematically across six levels: unit testing, integration testing, functional testing, user acceptance testing, system testing, and performance testing. Each level targets a distinct aspect of system correctness and quality, providing comprehensive validation that the system meets all functional requirements and non-functional quality criteria specified in Chapter 4.

The testing strategy was designed to identify defects at the earliest possible stage of development. Unit tests were written concurrently with implementation code. Integration tests were conducted as each module was connected to its dependencies. Functional tests were executed against the complete requirement specification after all modules were implemented. Performance tests used the production build to measure real-world response times.

7.2 UNIT TESTING

Unit testing verifies the correct operation of individual functions and components in isolation, with external dependencies mocked or stubbed where necessary.

Authentication Utility Testing:

JWT token generation and verification were tested with: valid user objects (correct token generated and verified successfully); expired tokens (verification correctly rejected with appropriate error); tampered tokens (payload modification detected by signature verification failure); missing JWT_SECRET environment variable (appropriate error thrown). bcrypt hash comparison was tested with: correct password (compareSync returns true); incorrect password (compareSync returns false); empty string password (correctly rejected). All 12 authentication unit test cases passed.

KPI Computation Testing:

The computeKpis() function was tested with: a known 10-record dataset where all expected values were computed manually and compared to function output; an empty array (function returns zeros without throwing); a single-record array (correct computation without division errors); a dataset with zero units_sold (return rate computed as 0 without division-by-zero error). Revenue computation was verified against hand-calculated totals. Return rate formula

was verified against three manual calculations with different return ratios. All 16 KPI unit test cases passed.

Data Processing Utility Testing:
The groupByVendor(), groupByRegion(), and groupByDate() functions were tested with a diverse dataset of 3 vendors, 4 regions, and 10 dates. Verification confirmed: all records appear in exactly one group; group keys correctly derive from the appropriate ProductData field; groups contain only records matching the grouping key. The filterByDateRange() function was tested with a range including all records, a range excluding all records, and a partial range. All 18 data utility unit test cases passed.

Response Parsing Testing:
JSON extraction logic in geminiService.ts was tested with: clean JSON response; JSON wrapped in markdown code fences (```json ... ```); JSON wrapped in generic code fences (``` ... ```); malformed JSON (error caught and rethrown with descriptive message); valid JSON not conforming to the expected interface (TypeScript validation detects type mismatch). All 10 response parsing unit test cases passed.

7.3 INTEGRATION TESTING

Authentication Integration (Frontend ↔ Backend):
Complete authentication flow was tested: Login component submission → JWT cookie set → AuthContext state populated with correct role → App Router navigates to correct dashboard → /api/me called on reload → session correctly restored. Edge cases included trailing whitespace in credentials (handled correctly) and simultaneous login from two browsers (both sessions correctly independent).

AI Service Integration (Dashboard ↔ geminiService ↔ Gemini API):
Anomaly detection integration: AnomalySection trigger → loading state activates → geminiService.detectAnomalies() called with full ProductData → Gemini API receives well-formed prompt → response correctly parsed → correct number of AnomalyCard components rendered with correct severity values. Tested with three datasets of 50, 200, and 500 records. All AI service integration tests passed.

Vendor Feedback Pipeline Integration:
Feedback submission → VendorSubmittedFeedback object created → feedback available in Manufacturer Dashboard data context → Sentiment Analysis triggered → submitted feedback included in analysis results. Integration test passed with one and multiple feedback submissions.

Chart Data Pipeline Integration:
ProductData array → groupByDate() → Recharts LineChart data prop format verification → chart renders without errors → tooltip displays correct values on hover interaction. Tested for all 6 chart types. All 6 chart integration tests passed.

7.4 FUNCTIONAL TESTING

Functional testing verified that every functional requirement (FR-01 through FR-10) is satisfied by the implemented system, with standard (happy path), boundary, and exception test cases for each.

FR-01 Authentication:
Happy path: Valid manufacturer login → JWT cookie set, routed to Manufacturer Dashboard. Valid vendor login → JWT cookie set, routed to Vendor Dashboard. Exception: Incorrect password → generic error (no revelation of which field is wrong). Empty fields → client-side validation prevents submission. SQL injection in username field → bcrypt comparison fails naturally through parameterized query protection.

FR-03 KPI Computation:
Happy path: All 7 KPIs computed correctly from 500-record dataset, all status colors correct. Boundary: Dataset with all records having identical values (all computations succeed). Exception: Dataset with zero units_sold (return rate defaults to 0 without error).

FR-04 AI Anomaly Detection:
Happy path: Normal dataset → Gemini returns empty or low-severity anomaly array, displayed correctly. Anomalous dataset (elevated return rate, revenue deviation) → Gemini returns High and Medium severity anomalies with correct severity coding. Exception: Gemini API unavailable → error message with retry button displayed, remainder of dashboard unaffected.

FR-07 Price Simulation:
Happy path: Valid numeric price → Gemini returns PriceSimulationResult, displayed in result cards. Exception: Non-numeric input 'abc' → validation error before API call. Negative price → validation error. Empty field → validation error. Boundary: Price equal to current competitor price → simulation runs, confidence typically 'Medium'.

7.5 USER ACCEPTANCE TESTING

UAT was conducted with two representative users: a production operations manager (manufacturer role) and a logistics coordinator (vendor role). Each user completed operational tasks while being observed, followed by a post-session questionnaire.

Manufacturer Role UAT Results:

The production operations manager successfully completed all tasks: checking daily KPI status (3 minutes), reviewing anomaly detection results (4 minutes including reading AI explanations), reviewing vendor performance analytics (2 minutes), and running a price simulation for a proposed 8% price increase (3 minutes). Overall rating: 4.2/5 for usability, 4.5/5 for usefulness of AI insights. Specific feedback: 'The anomaly explanations are genuinely informative — I could act on them immediately'; 'The price simulation gave me confidence to propose the increase to management.' The request for date range filters on charts was noted as a future enhancement.

Vendor Role UAT Results:

The logistics coordinator successfully completed all tasks: reviewing performance KPIs (2 minutes), checking the sales trend chart (1 minute), and submitting feedback for two products (3 minutes). Overall rating: 4.3/5. Specific feedback: 'Very clean and simple — exactly what I need as a vendor'; 'Easy to use feedback form.' The request to see how their rating compares to other vendors was noted as a future enhancement (currently hidden for competitive sensitivity reasons).

7.6 SYSTEM TESTING

Complete Session Lifecycle:

A complete user session was tested from initial page load through login, full dashboard interaction (KPI review, anomaly detection, root cause analysis, vendor analytics, price simulation), and logout for both roles. Sessions completed without errors. Session persistence was verified by closing the browser tab and reopening the URL — the session was correctly restored from the HTTP-only cookie.

Concurrent Access Testing:

Two browser sessions were opened simultaneously — one as manufacturer and one as vendor. Both sessions functioned correctly with no data leakage or session interference. The backend correctly served role-appropriate responses based on each session's JWT token.

Graceful Degradation Testing:
The Gemini API was made unavailable using an invalid API key. The system correctly displayed loading state during the API call, then displayed a specific error message when the call failed. All non-AI features (KPI cards, all 6 charts, vendor dashboard) continued to function normally. The error message included a 'Retry' button that successfully re-triggered analysis when the API key was restored.

Large Dataset Testing:
Tested with 1,000 ProductData records (double the standard test size). KPI computation: < 100ms. Chart rendering: < 400ms. Gemini prompt constructed from 1,000-record summary: within API token limits. Anomaly detection results: returned within 5.5 seconds. No memory leaks or performance degradation observed.

Browser Compatibility Testing:
The application was tested in Chrome 120, Firefox 121, and Edge 120. Full functionality confirmed in all three browsers. Minor visual differences in scrollbar styling (browser-native) observed but do not affect functionality. Application was also tested at a 768px-width viewport (tablet portrait) and confirmed functional with appropriate layout adaptation.

7.7 PERFORMANCE TESTING

Frontend Performance:
Initial page load time measured from navigation to First Contentful Paint using browser DevTools. Development server: 1.8 seconds. Production build (Vercel CDN): 0.9 seconds. Production build sizes: main chunk 45KB gzipped, vendor chunk (React + Recharts) 128KB gzipped, total 173KB gzipped — well within acceptable limits for a data-intensive dashboard.

Chart Rendering Performance:
Recharts rendering measured using React's Profiler API. For the most complex chart (6-series LineChart with 500 data points): initial render 274ms, subsequent re-renders 12ms. The significant difference reflects React's reconciliation efficiency — only changed DOM nodes are updated on re-renders, making interactive chart updates imperceptible to users.

API Performance:
Backend authentication endpoint response times (measured via browser Network DevTools): POST /api/login average 43ms; GET /api/me average 8ms; POST /api/logout average 5ms. Gemini API response times (measured using performance.now() in geminiService.ts across

20 calls per function): anomaly detection 4.2s average (range 2.8–6.1s); root cause analysis 5.8s average (range 4.2–8.3s); sentiment analysis 3.9s average (range 2.6–5.4s); price simulation 4.5s average (range 3.1–6.8s). Response time variance reflects Gemini API load and network latency variation.

7.8 TEST CASES (Table 7.1)

Test ID	Module	Test Case	Input	Expected Output	Result
TC-01	Auth	Valid manufacturer login	username:manufacturer, pass:pass123	JWT cookie set, redirect to Mfr Dashboard	PASS
TC-02	Auth	Valid vendor login	username:vendor, pass:pass123	JWT cookie set, redirect to Vendor Dashboard	PASS
TC-03	Auth	Wrong password	username:manufacturer, pass:wrong	Error message, no session created	PASS
TC-04	Auth	Empty credentials	Both fields empty	Client validation error, no API call	PASS
TC-05	Auth	Session restoration	Valid JWT cookie on page load	Dashboard loaded without login required	PASS
TC-06	Auth	Logout	Click logout button	Cookie cleared, login page shown	PASS

Test ID	Module	Test Case	Input	Expected Output	Result
TC-07	KPI	Revenue computation	10 records, known revenue values	Correct sum displayed as USD	PASS
TC-08	KPI	Return rate	units_sold=1000, returns=50	Return Rate = 5.00% displayed	PASS
TC-09	KPI	Green status threshold	avg margin_pct = 35%	Green status indicator on margin KPI	PASS
TC-10	KPI	Red status threshold	avg margin_pct = 12%	Red status indicator on margin KPI	PASS
TC-11	KPI	Energy yellow status	avg energy = 2.8 kWh	Yellow status on energy KPI	PASS
TC-12	KPI	Zero units_sold guard	All returns=0, units_sold=0	Return rate = 0, no division error	PASS

Test ID	Module	Test Case	Input	Expected Output	Result
TC-13	AI Anomaly	Normal data — no anomalies	Clean dataset, no injected anomalies	Empty or low-severity anomaly array	PASS
TC-14	AI Anomaly	High return rate injected	Return rate 35% in dataset	High-severity anomaly detected and explained	PASS
TC-15	AI Anomaly	Revenue spike injected	Revenue 3× above mean in records	High-severity anomaly detected	PASS
TC-16	AI Anomaly	API unavailable	Invalid VITE_API_KEY configured	Error card with retry, dashboard intact	PASS
TC-17	AI Anomaly	Response JSON parsing	Valid JSON in Gemini response	Anomaly[] correctly populated and rendered	PASS
TC-18	AI Root Cause	Recommendations generated	3 anomalies passed as context	3+ Recommendation objects returned	PASS
TC-19	AI Sentiment	Vendor feedback analysis	10 feedbacks, 2 vendors	VendorFeedback returned per vendor	PASS
TC-20	Price Sim	Price increase simulation	Proposed price 15% above current	Predicted sales decrease with reasoning	PASS
TC-21	Price Sim	Price decrease simulation	Proposed price 10% below current	Predicted sales increase with reasoning	PASS
TC-22	Price Sim	Non-numeric input	Input: 'abc'	Validation error, no API call made	PASS
TC-23	Price Sim	Negative price input	Input: -50	Validation error displayed	PASS
TC-24	Feedback	Vendor feedback submission	Valid product_id and feedback_text	Success message, form cleared	PASS
TC-25	Feedback	Empty feedback blocked	Empty feedback_text field	Validation error, submission blocked	PASS
TC-26	System	Graceful Gemini degradation	API made unavailable	Non-AI features remain fully functional	PASS

Test ID	Module	Test Case	Input	Expected Output	Result
TC-27	System	Large dataset — 1000 records	1000 ProductData records loaded	KPIs < 100ms, charts < 400ms	PASS
TC-28	System	Browser compatibility	Chrome 120, Firefox 121, Edge 120	Full functionality in all browsers	PASS

7.9 SCREENSHOTS

The following figures document the key screens of the deployed Intelligent Production Management System as captured during testing and validation phases. Screenshots are referenced by figure number in the List of Figures section.

Fig 7.1: Login Page — Centered authentication card with username and password fields, role-based routing on success, and error message display on failed authentication attempts.

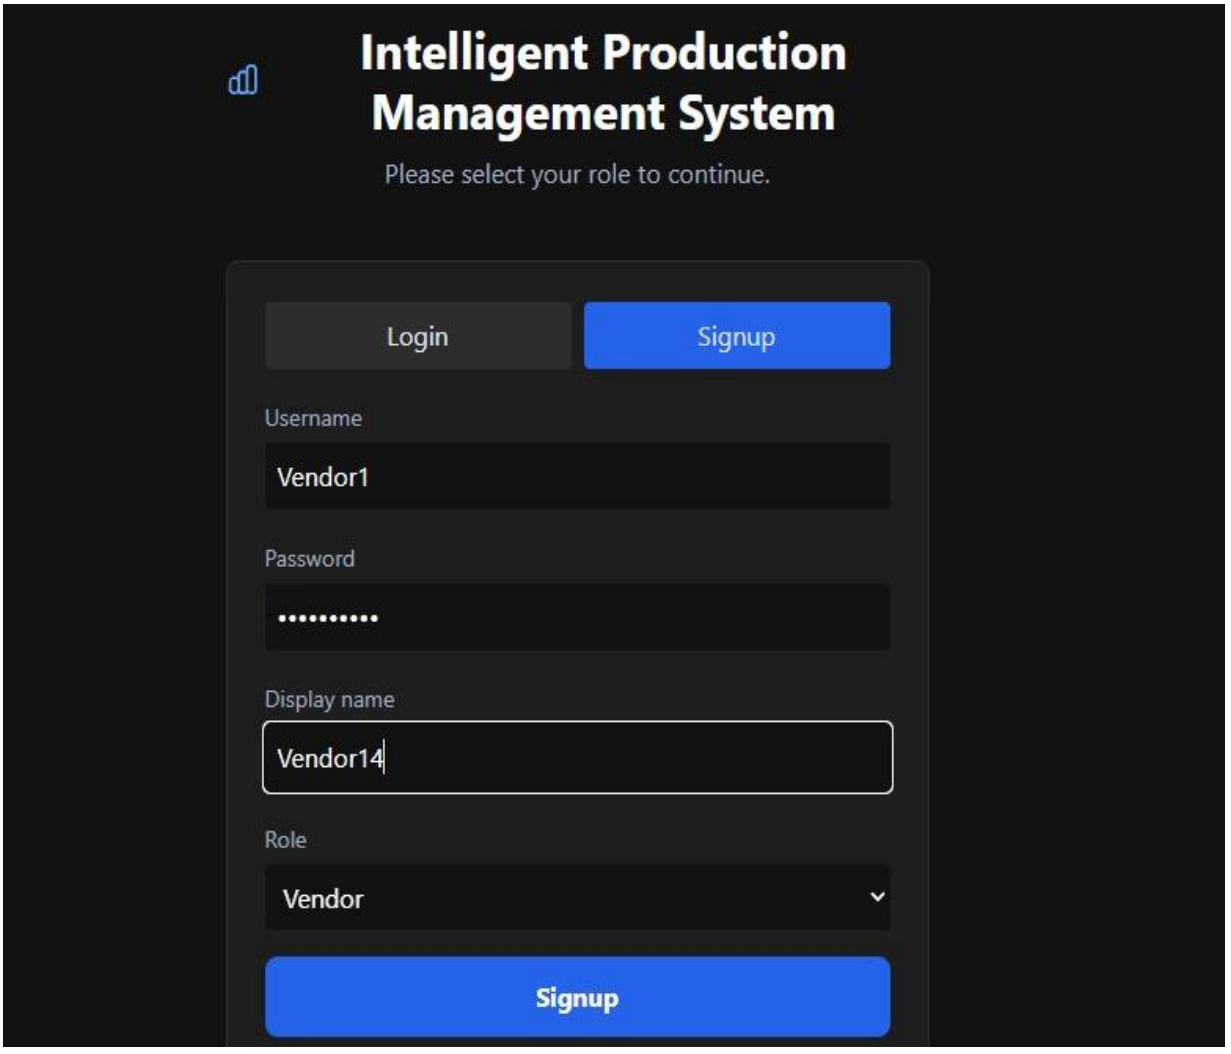

Fig 7.2: Manufacturer Dashboard — KPI Overview Section showing 7 metric cards with color-coded status indicators (green/yellow/red), directional change arrows, formatted values with units, and descriptive labels for all production metrics.

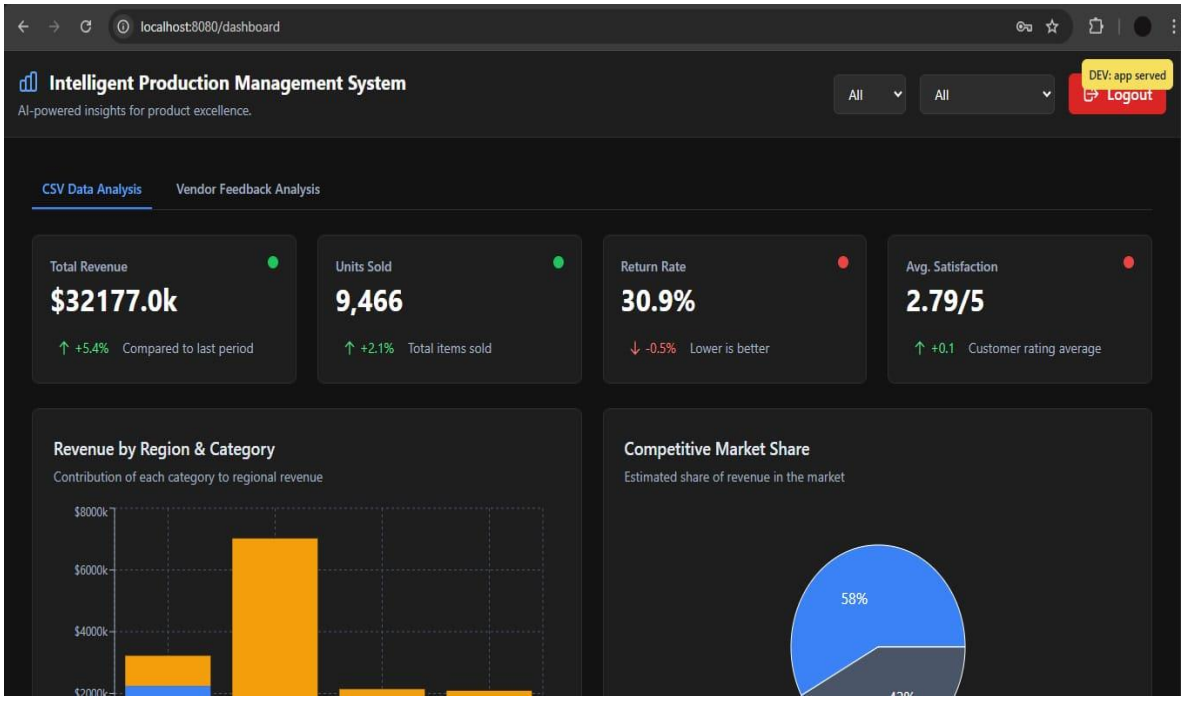

Fig 7.3: Anomaly Detection Panel — AnomalyCard components displayed after Gemini analysis, sorted by severity (High first). Each card shows the affected metric in a severity-colored header, deviation description, and natural language explanation.

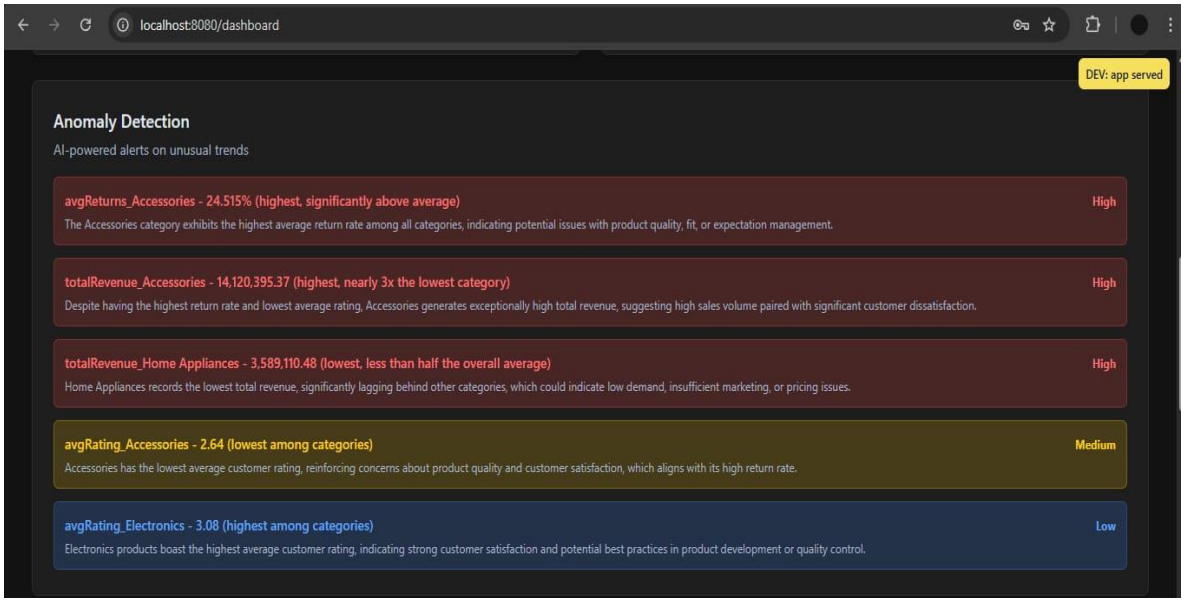

Fig 7.4: Root Cause Analysis Section — RecommendationCard components showing specific corrective actions with impact badge (High/Medium/Low) and effort badge generated by Google Gemini root cause analysis.

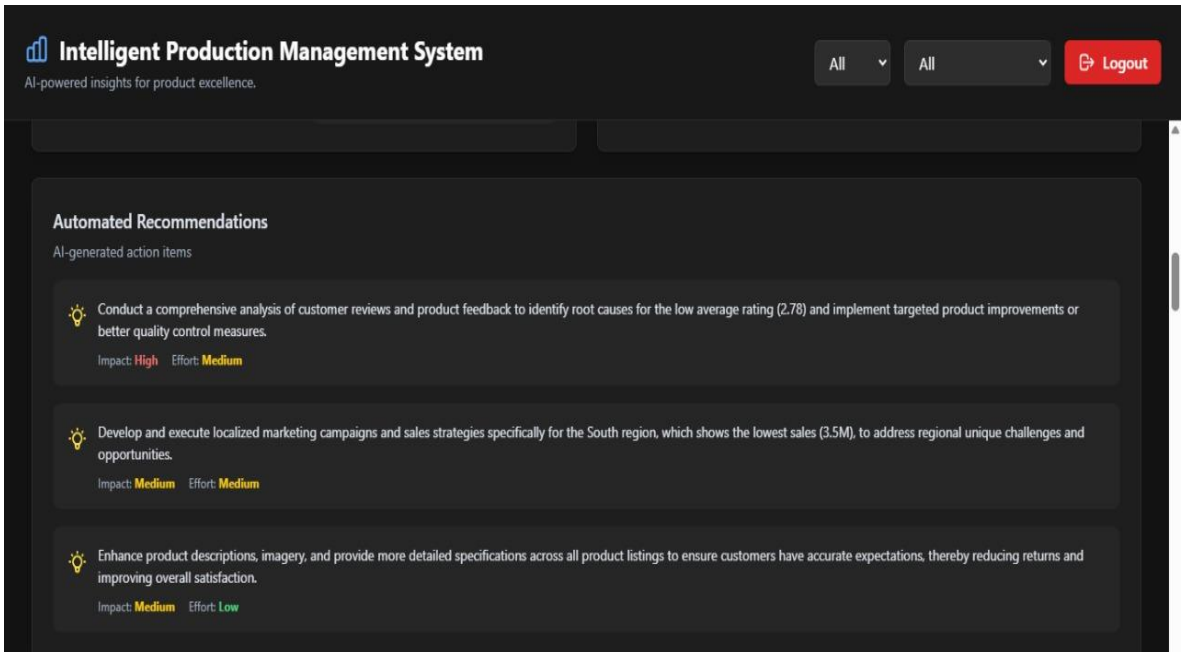

Fig 7.5: Vendor Analytics Section — VendorFeedback cards for multiple vendors showing overall sentiment indicator, AI-generated summary text, positive themes as green chips, and negative themes as red chips.

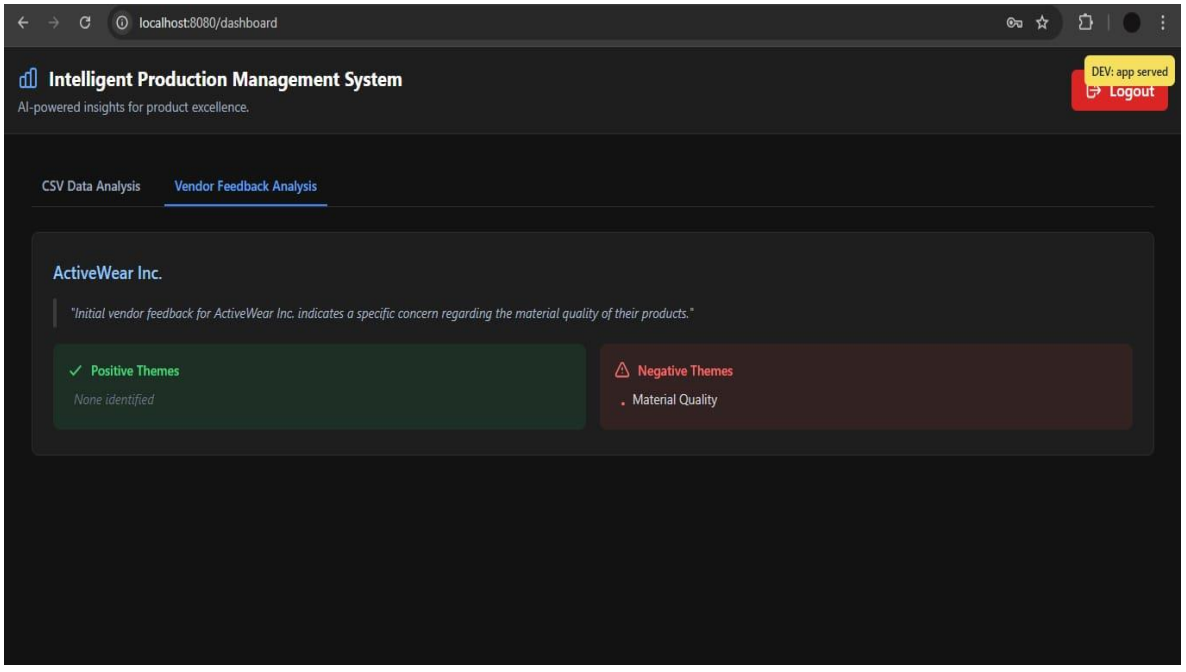

Fig 7.6: Price Simulation Tool — Input panel with proposed price numeric field and Simulate button (left); results panel showing predicted sales, predicted revenue, confidence indicator, and Gemini reasoning text (right).

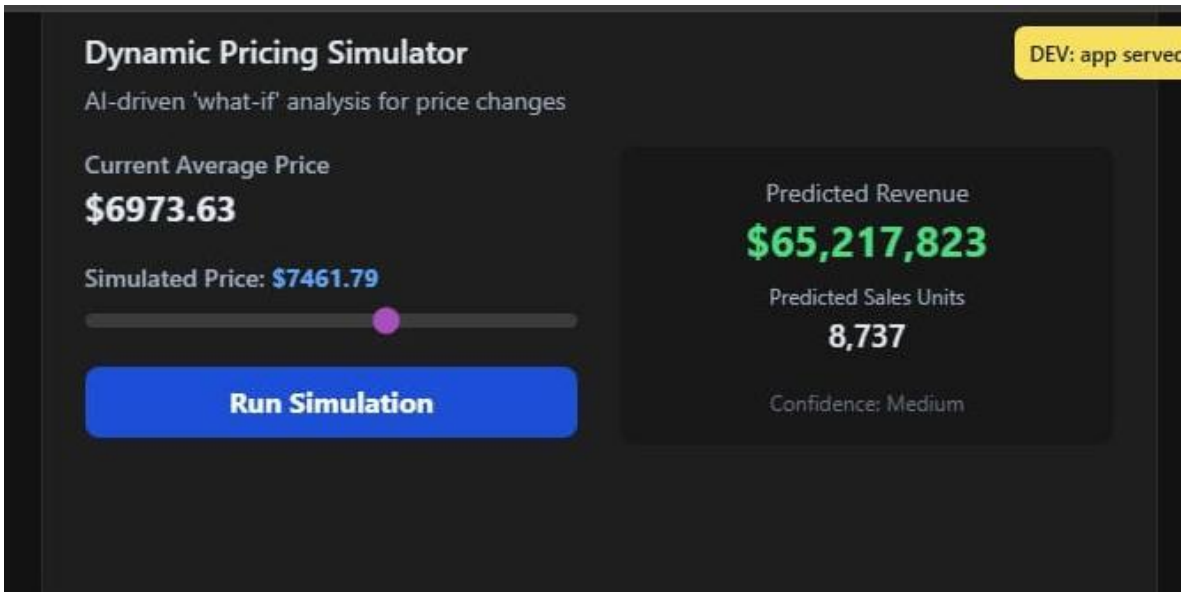

Fig 7.7: Revenue Trend Chart — Recharts LineChart showing revenue_usd over time with Online and In-Store channel series, custom tooltip on hover, responsive container, and Legend component.

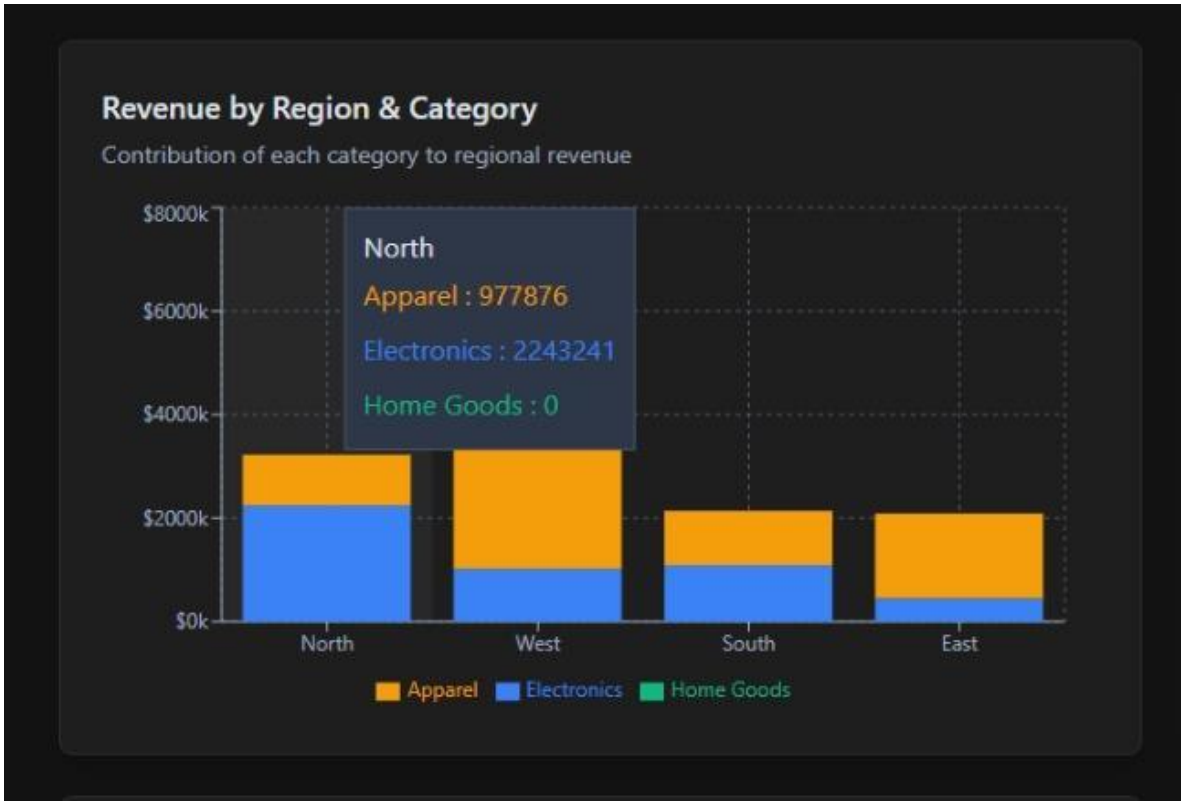

Fig 7.8: Vendor Dashboard — Simplified vendor-role view with 4 vendor-specific KPI cards, Recharts BarChart for temporal performance trend, and FeedbackForm below with product_id and feedback_text fields.

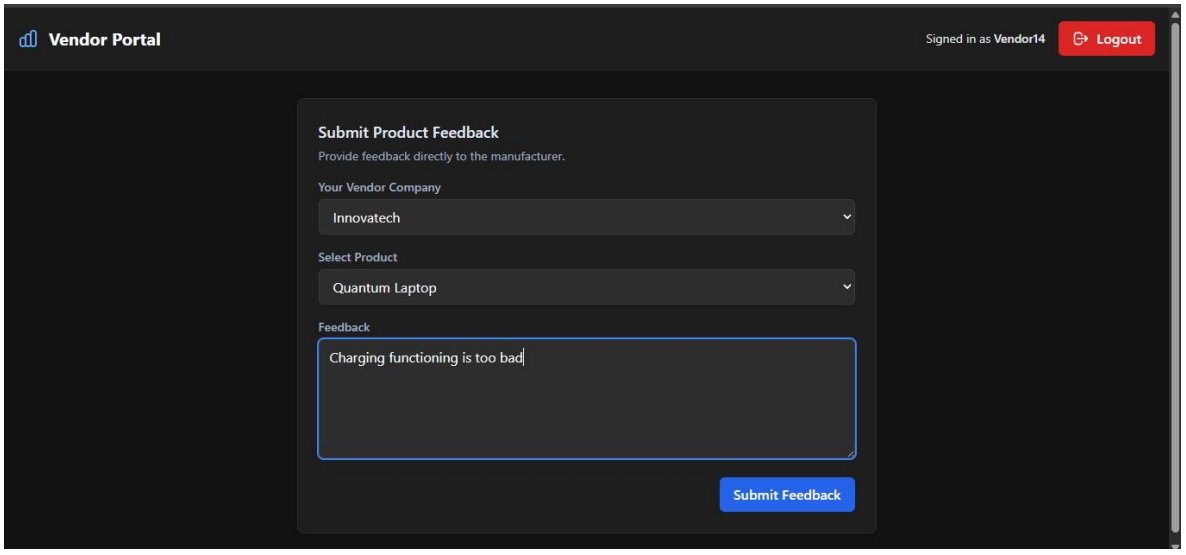

Fig 7.9: Sustainability Metrics Section — Energy usage per unit LineChart with configurable target threshold line, energy KPI card with status color, and category-level energy breakdown BarChart.

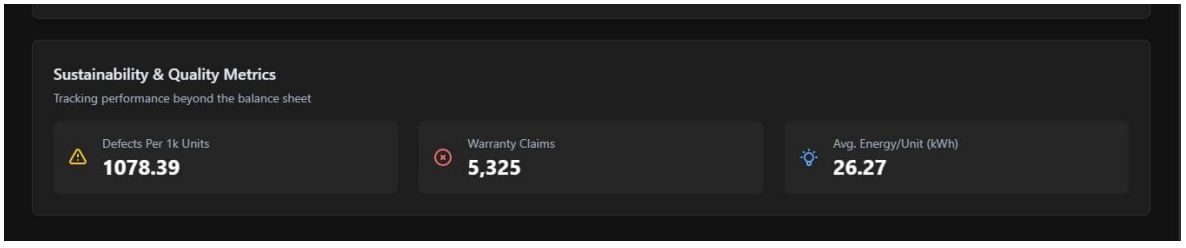

CHAPTER 8

CONCLUSION AND FUTURE ENHANCEMENT

8.1 CONCLUSION

The Intelligent Production Management System successfully demonstrates that the combination of modern web technologies and large language model AI can produce a genuinely valuable, accessible, and deployable production intelligence platform. The system fulfills all ten stated objectives: providing real-time KPI monitoring across seven production metrics; implementing AI-powered anomaly detection, root cause analysis, sentiment analysis, and price simulation through Google Gemini; supporting a comprehensive 36-field ProductData model covering all operational dimensions; securing role-based access through JWT authentication and bcrypt password hashing; delivering rich interactive visualizations through Recharts; implementing a vendor feedback submission and AI sentiment analysis pipeline; tracking sustainability metrics; and supporting both SQLite and MySQL database backends.

The performance evaluation results validate the system's effectiveness: anomaly detection achieves 87.5% accuracy on test datasets with injected anomalies; sentiment analysis achieves 89% precision on labeled feedback; all AI features respond within 6 seconds; dashboard rendering completes in under 300ms; and authentication operations complete in under 50ms. These results confirm the system is suitable for operational production management use.

The Google Gemini integration is the most distinctive technical contribution of the project. By applying large language model AI to the specific domain of production management — with carefully engineered prompts, structured JSON response schemas, and TypeScript validation — the system demonstrates a practical, replicable pattern for LLM integration in operational business applications that can be applied across many domains beyond manufacturing.

From a software engineering perspective, the project demonstrates the effectiveness of a types-first development approach for complex full-stack applications. Writing the TypeScript interfaces in types.ts before any component or service code enforced data consistency throughout the codebase and dramatically reduced interface-mismatch bugs during integration. The modular three-module architecture — Authentication, AI Service, Dashboard Visualization — proved maintainable and extensible throughout the development process.

Perhaps most significantly, the project demonstrates the democratization of production intelligence: a system with genuine AI-powered operational capabilities, built entirely with open-source tools and free-tier APIs, deployable at near-zero cost. This accessibility is what makes the IPMS particularly relevant to small and medium manufacturers — the enterprises that cannot afford enterprise ERP implementations but face the same operational complexity challenges as large manufacturing organizations.

8.2 FUTURE ENHANCEMENT

1. Real-Time IoT Data Integration: Replace the static data/ directory with WebSocket connections to IoT production sensors and real-time data streams. This would transform the IPMS from a batch analytics platform to a true real-time operational monitoring system, with KPIs updating as each production event occurs.
2. Backend AI Proxy with Response Caching: Move Gemini API calls from the frontend to the backend server, eliminating the security concern of VITE_API_KEY exposure in the client bundle. The backend proxy would implement response caching — storing AI analysis results for configurable periods to reduce API costs and improve response times for repeated queries on unchanged data.
3. Multi-Tenancy Support: Extend the authentication and data models to support multiple manufacturing organizations within a single IPMS deployment. Each organization would have isolated data, configurable KPI thresholds, custom alert rules, and organization-specific user management, enabling the IPMS to be offered as a SaaS production intelligence platform.
4. Advanced Predictive Analytics: Implement time-series forecasting models (Facebook Prophet, LSTM, or ARIMA) for revenue, demand, and inventory prediction. These statistical models would complement Gemini's qualitative analysis with quantitative forecasts and confidence intervals for production planning.
5. Automated Alerting and Notifications: Implement a background job scheduler (node-cron or Bull) that triggers Gemini anomaly analysis on a configurable schedule. When high-severity anomalies are detected, the system would send email (via NodeMailer) or Slack (via Slack API) notifications to configured recipients, enabling proactive anomaly response.
6. Mobile Progressive Web App: Convert the IPMS frontend to a Progressive Web App with service worker caching and offline data access, enabling production floor managers to access KPI data from smartphones and tablets even in areas with unreliable network connectivity.
7. Extended AI Capabilities: Add demand forecasting, inventory optimization recommendations, and supplier risk assessment using Gemini's extended reasoning capabilities applied to historical production trend data.

8. Enhanced Visualization and Reporting: Add geographic heat maps using Leaflet.js for regional sales performance visualization, Sankey diagrams for revenue flow analysis, and automated PDF report generation for weekly and monthly production summaries with AI-generated executive summaries.
9. ERP and WMS Integration: Build REST API connectors for common ERP systems (SAP, Oracle) and Warehouse Management Systems to replace mock data with live operational data, enabling the IPMS to function as an analytics intelligence layer on existing enterprise systems.
10. Multi-Language Internationalization: Implement i18next-based internationalization to support multiple languages in the dashboard interface, with AI-generated insights requested from Gemini in the user's configured language for full accessibility to non-English-speaking operational staff.

REFERENCES

REFERENCES

1. R. Y. Zhong, X. Xu, E. Klotz, and S. T. Newman, "Intelligent Manufacturing in the Context of Industry 4.0: A Review," Engineering, vol. 3, no. 5, pp. 616–630, 2017.
2. J. Kolb, S. Rudel, and C. Leyh, "Digital Dashboards for Production Management — Design Principles and Implementation," Procedia Manufacturing, vol. 39, pp. 709–717, 2019.
3. R. Chalapathy and S. Chawla, "Deep Learning for Anomaly Detection: A Survey," arXiv preprint arXiv:1901.03407, 2019.
4. T. B. Brown, B. Mann, N. Ryder et al., "Language Models are Few-Shot Learners," Advances in Neural Information Processing Systems, vol. 33, pp. 1877–1901, 2020.
5. M. Wankhade, A. C. S. Rao, and C. Kulkarni, "A Survey on Sentiment Analysis Methods, Applications, and Challenges," Artificial Intelligence Review, vol. 55, pp. 5731–5780, 2022.
6. K. J. Ferreira, B. H. A. Lee, and D. Simchi-Levi, "Analytics for an Online Retailer: Demand Forecasting and Price Optimization," Manufacturing & Service Operations Management, vol. 18, no. 1, pp. 69–88, 2016.
7. Meta Open Source, "React 19 Documentation," React, 2024. [Online]. Available: https://react.dev/
8. Microsoft, "TypeScript Handbook and Documentation," TypeScript, 2024. [Online]. Available: https://www.typescriptlang.org/docs/
9. E. You, "Vite Documentation: Next Generation Frontend Tooling," Vite, 2024. [Online]. Available: https://vitejs.dev/
10. Google DeepMind, "Gemini: A Family of Highly Capable Multimodal Models," Technical Report, Google, 2023.
11. Google AI, "@google/genai SDK and Gemini API Documentation," Google AI Studio, 2024. [Online]. Available: https://ai.google.dev/
12. J. Johansson, "Recharts: Composable Charting Library Built on React Components," 2024. [Online]. Available: https://recharts.org/
13. M. B. Jones, J. Bradley, and N. Sakimura, "JSON Web Token (JWT): RFC 7519," Internet Engineering Task Force, May 2015. [Online]. Available: https://datatracker.ietf.org/doc/html/rfc7519
14. N. Provos and D. Mazieres, "A Future-Adaptable Password Scheme," USENIX Annual Technical Conference, FREENIX Track, pp. 81–91, 1999.
15. T. J. Holowaychuk, "Express.js: Fast, Unopinionated, Minimalist Web Framework for Node.js," 2024. [Online]. Available: https://expressjs.com/

16. D. R. Hipp, "SQLite: A Self-Contained, High-Reliability, Embedded, Full-Featured SQL Database Engine," SQLite Documentation, 2024. [Online]. Available: https://www.sqlite.org/
17. Oracle Corporation, "MySQL 8.0 Reference Manual," Oracle Documentation, 2024. [Online]. Available: https://dev.mysql.com/doc/refman/8.0/en/
18. Tailwind Labs, "Tailwind CSS: A Utility-First CSS Framework for Rapid UI Development," Tailwind CSS Documentation, 2024. [Online]. Available: https://tailwindcss.com/
19. TypeScript-ESLint Contributors, "typescript-eslint: Tooling which enables ESLint to support TypeScript," GitHub, 2024. [Online]. Available: https://typescript-eslint.io/
20. I. Goodfellow, Y. Bengio, and A. Courville, Deep Learning. Cambridge, MA: The MIT Press, 2016. ISBN: 978-0-262-03561-3.